
Computer Graphics

Let there be Light

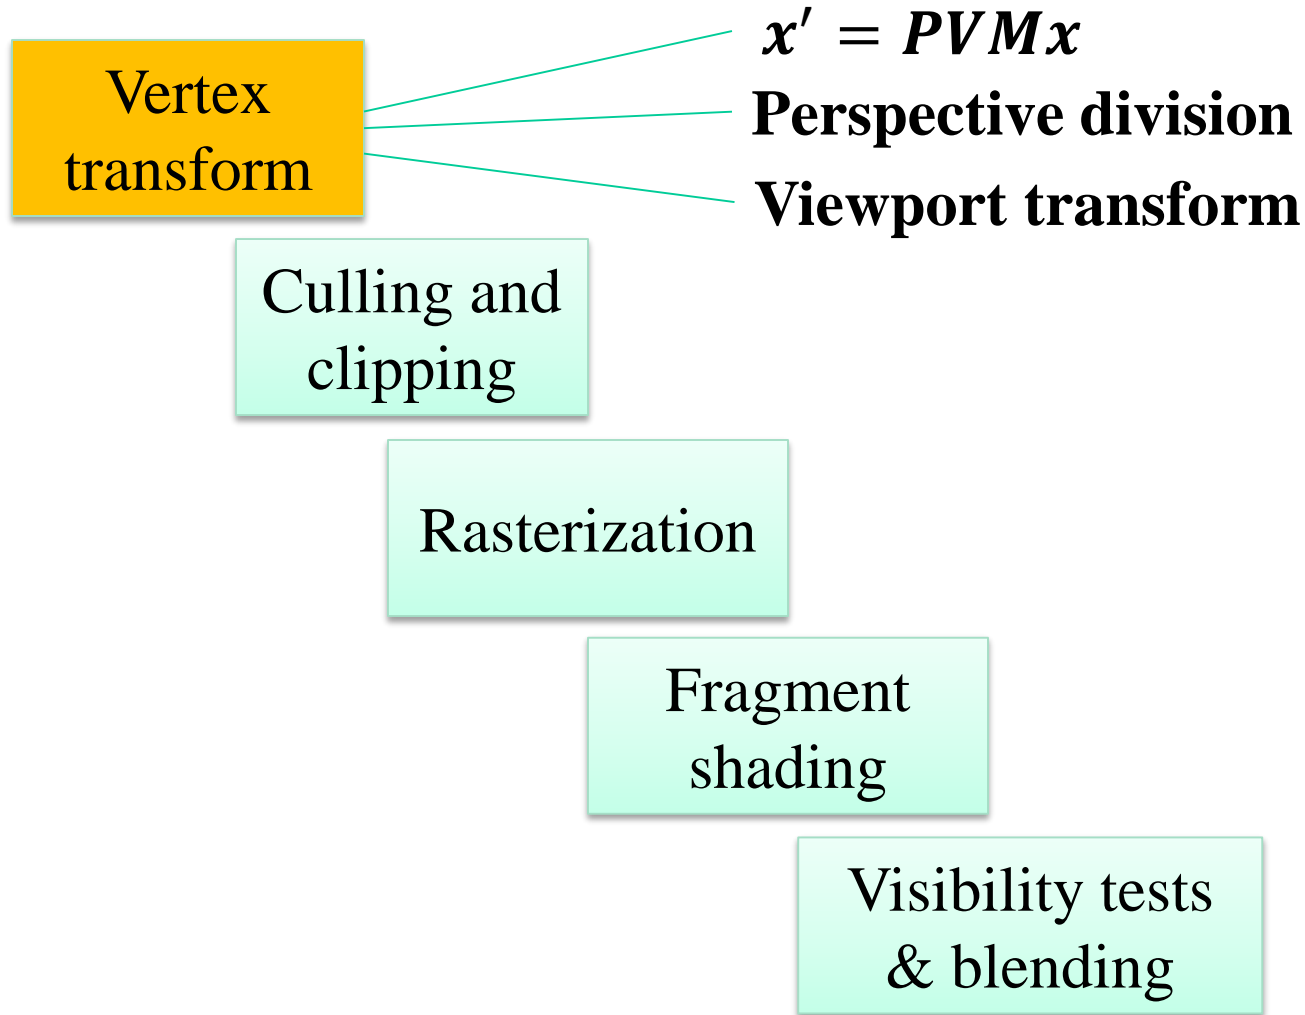
Konstantin Tretyakov

kt@ut.ee

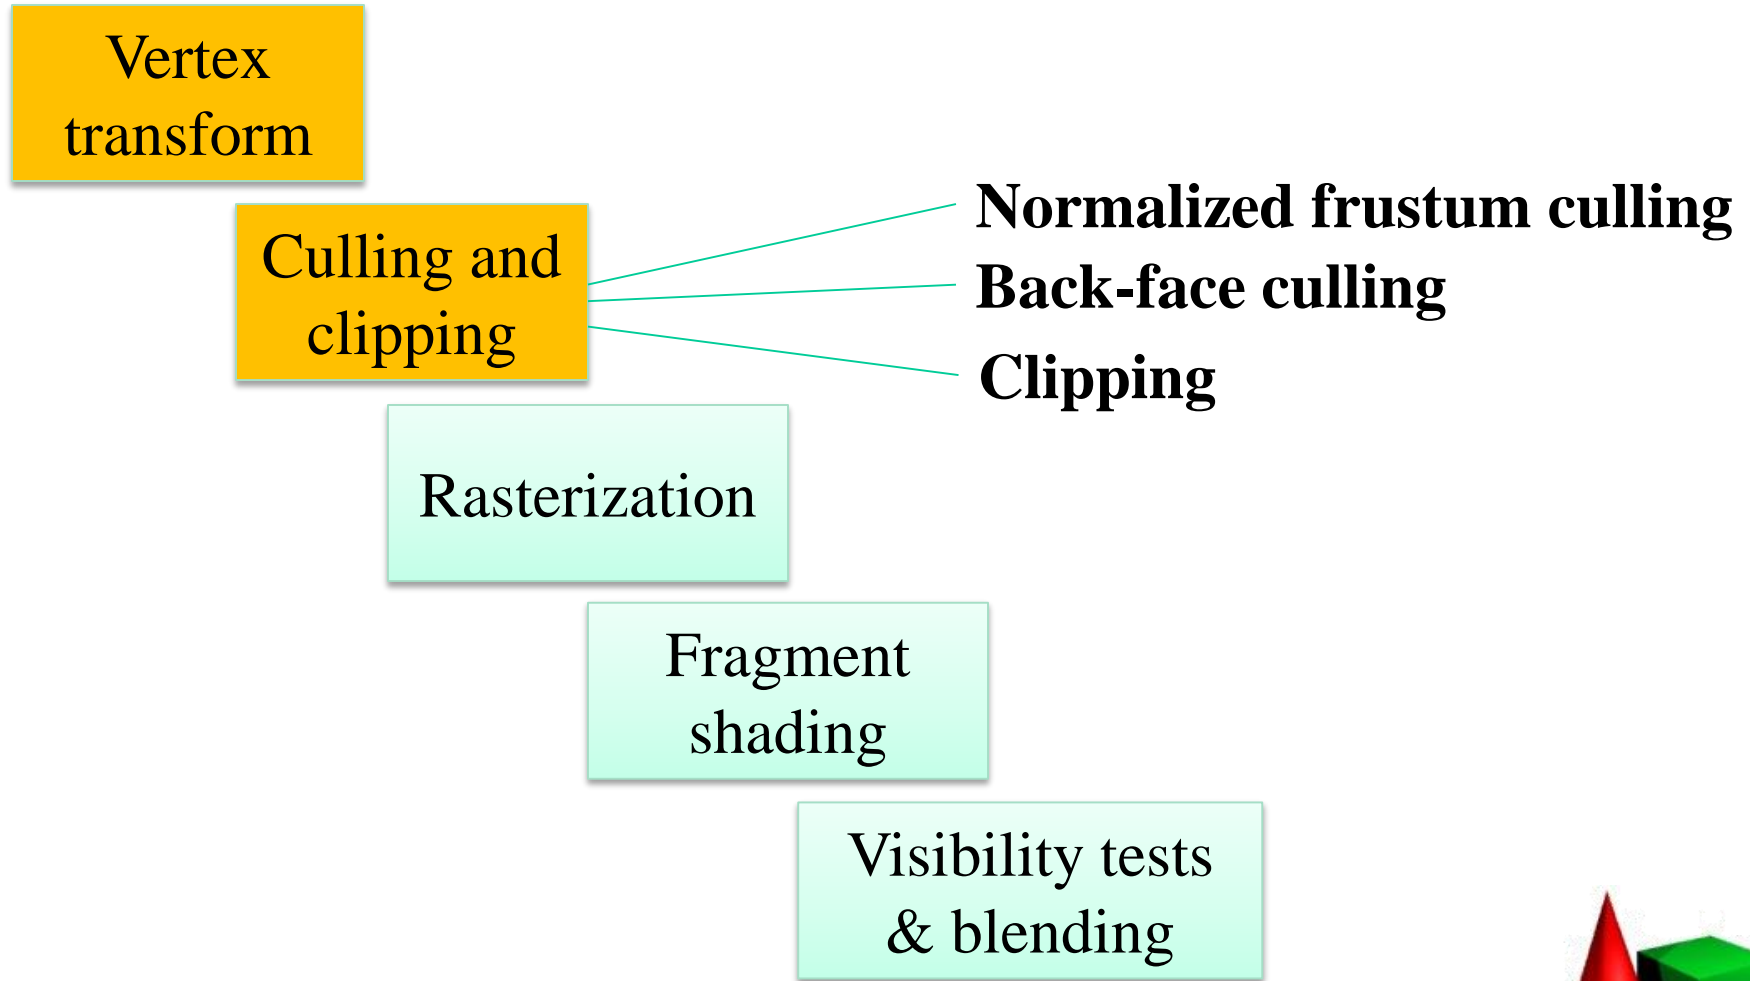


Oct 16, 2013

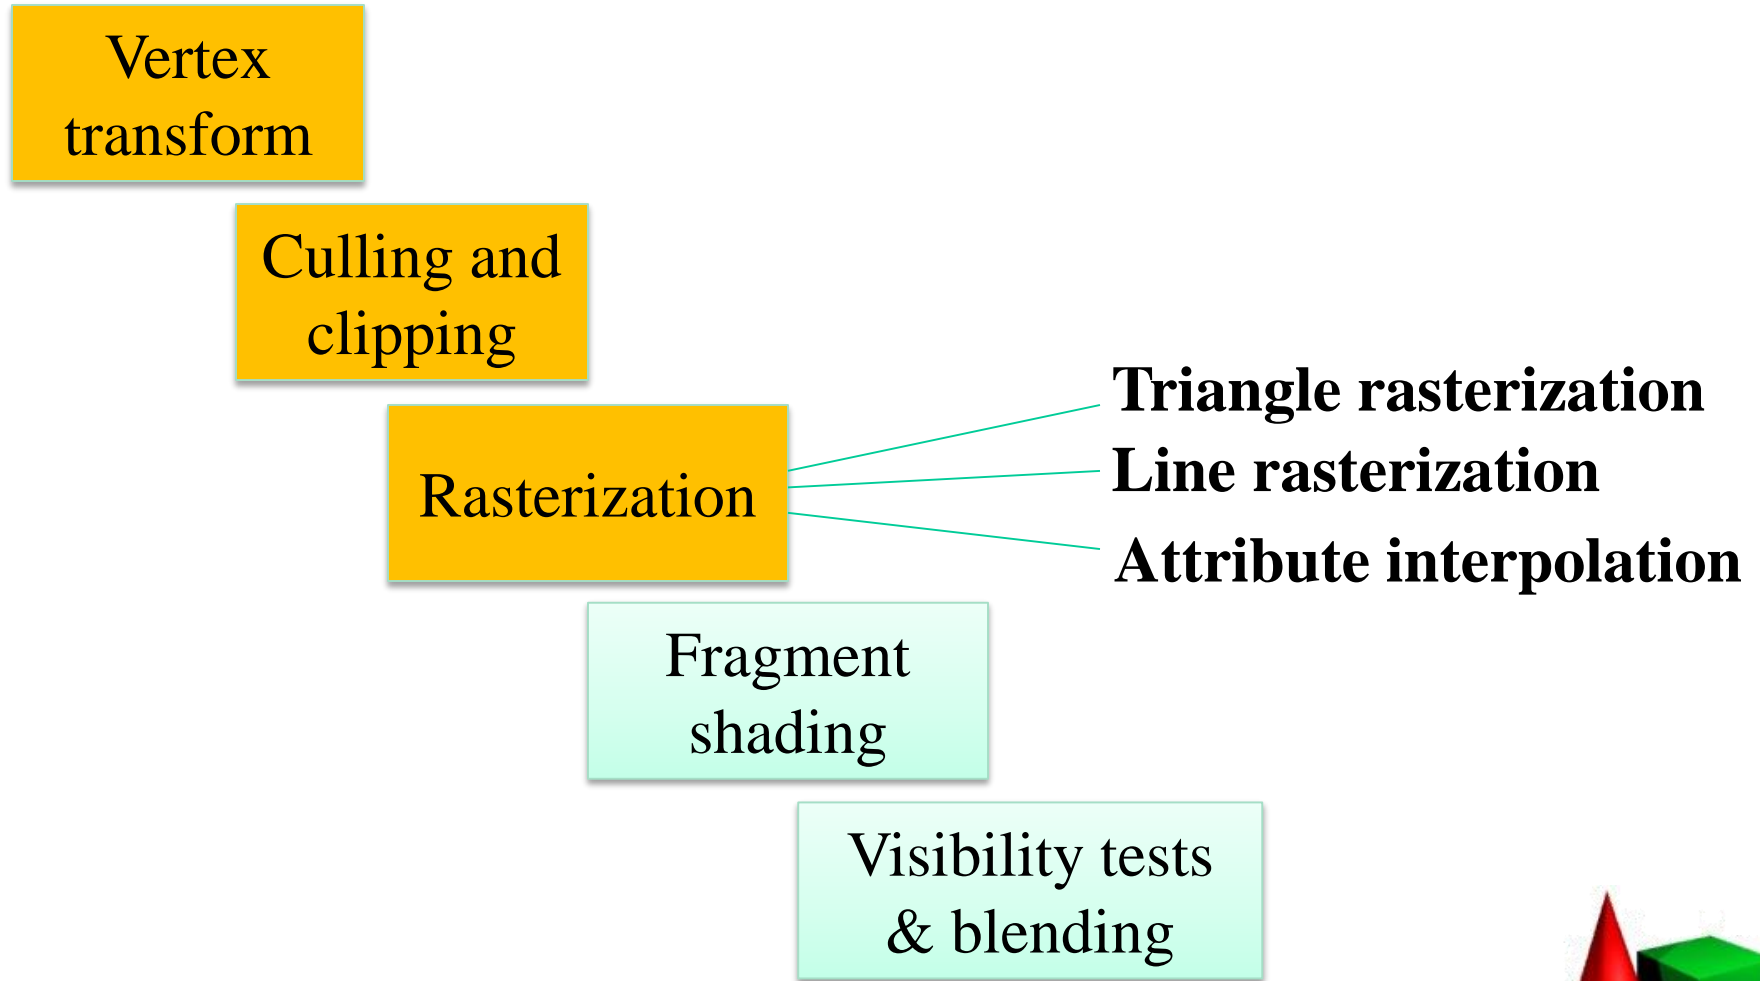
Standard Graphics Pipeline



Standard Graphics Pipeline



Standard Graphics Pipeline



Standard Graphics Pipeline

Vertex
transform

Culling and
clipping

Rasterization

Fragment
shading

Visibility tests
& blending

Blending
Z-buffer
Stencil buffer



Standard Graphics Pipeline

Vertex
transform

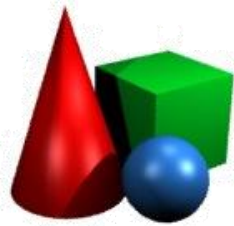
Culling and
clipping

Rasterization

Fragment
shading

Next

Visibility tests
& blending



Standard Graphics Pipeline

Vertex
transform

Determine clip-space position of a triangle

Culling and
clipping

Determine whether the triangle is visible

Rasterization

Determine all pixels belonging to the triangle

Fragment
shading

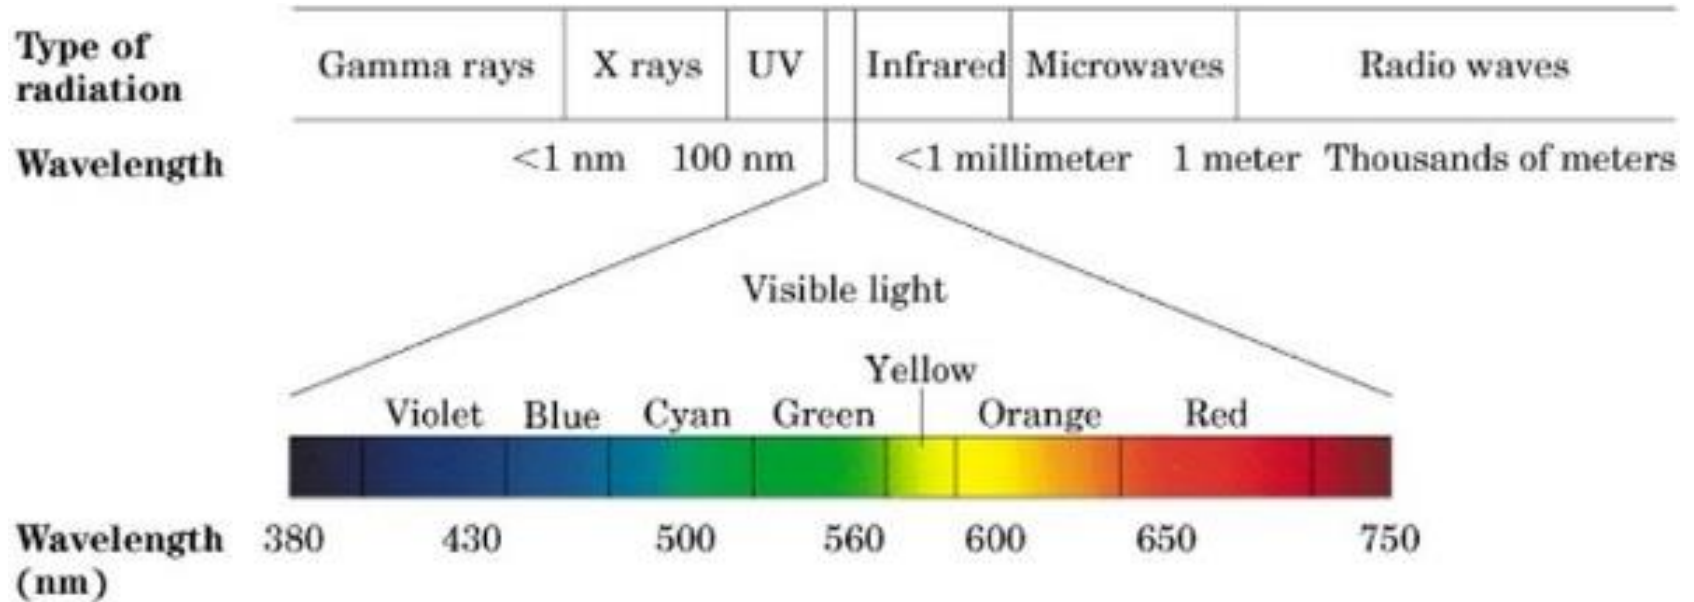
For each pixel, determine its **color**

Visibility tests
& blending

Draw pixel (*if needed*)



Light vs. Color

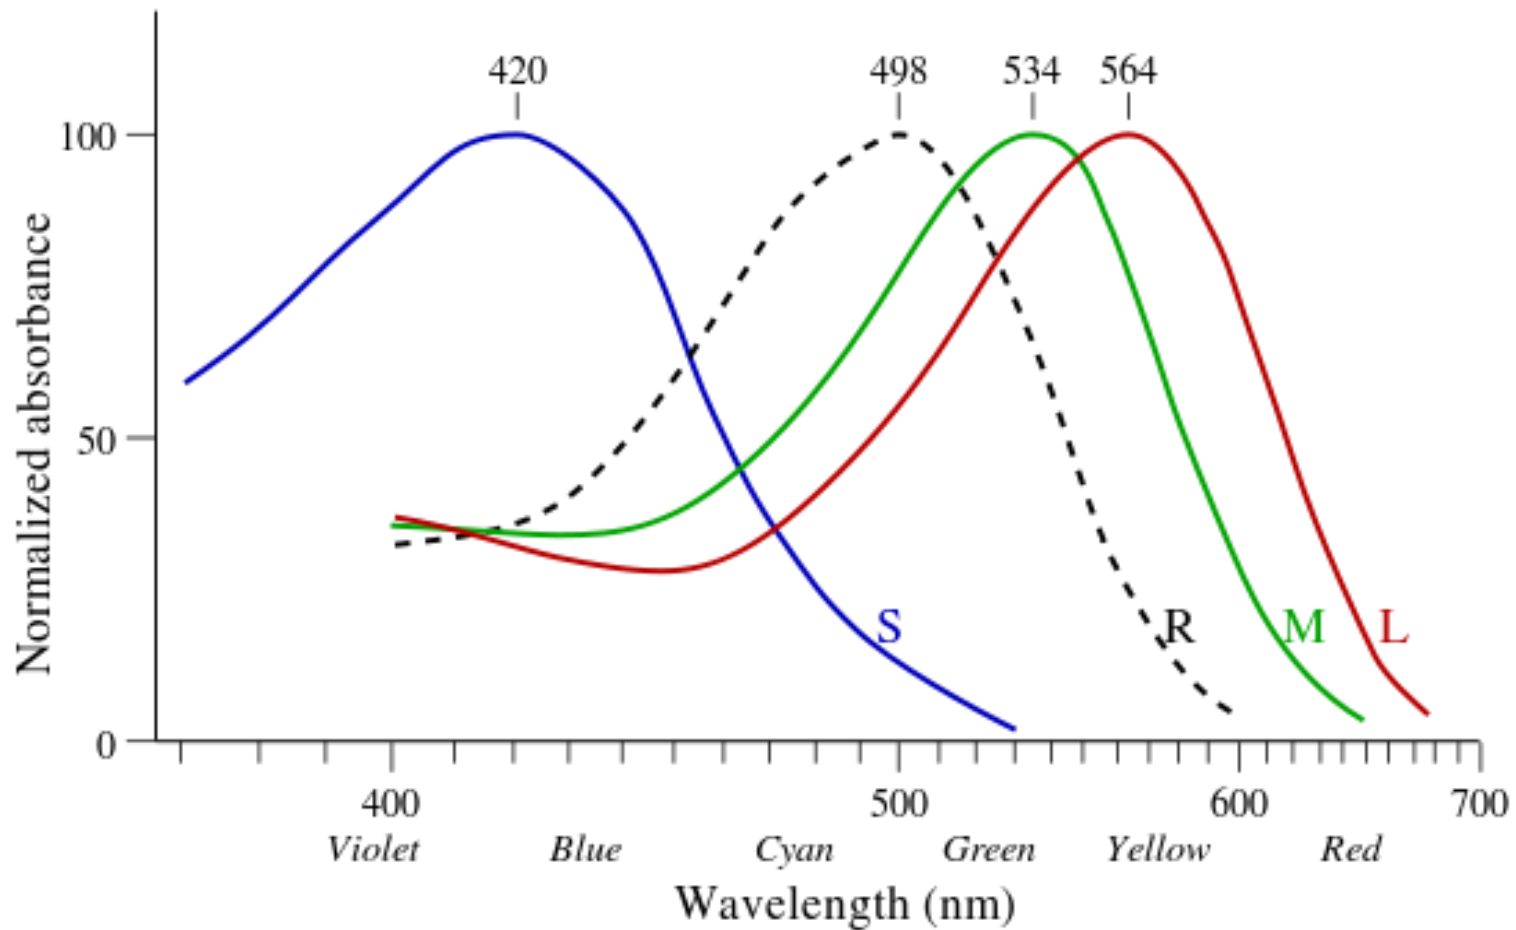


Light vs. Color

- In principle, a light wave can have a very complex spectrum (e.g. think how complex a sound wave can be).
- However, the receptor cells in our eyes only can only extract crude aggregates of this complex signal.



Rods & Cones



Rods & Cones

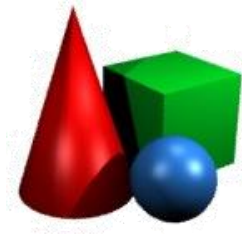
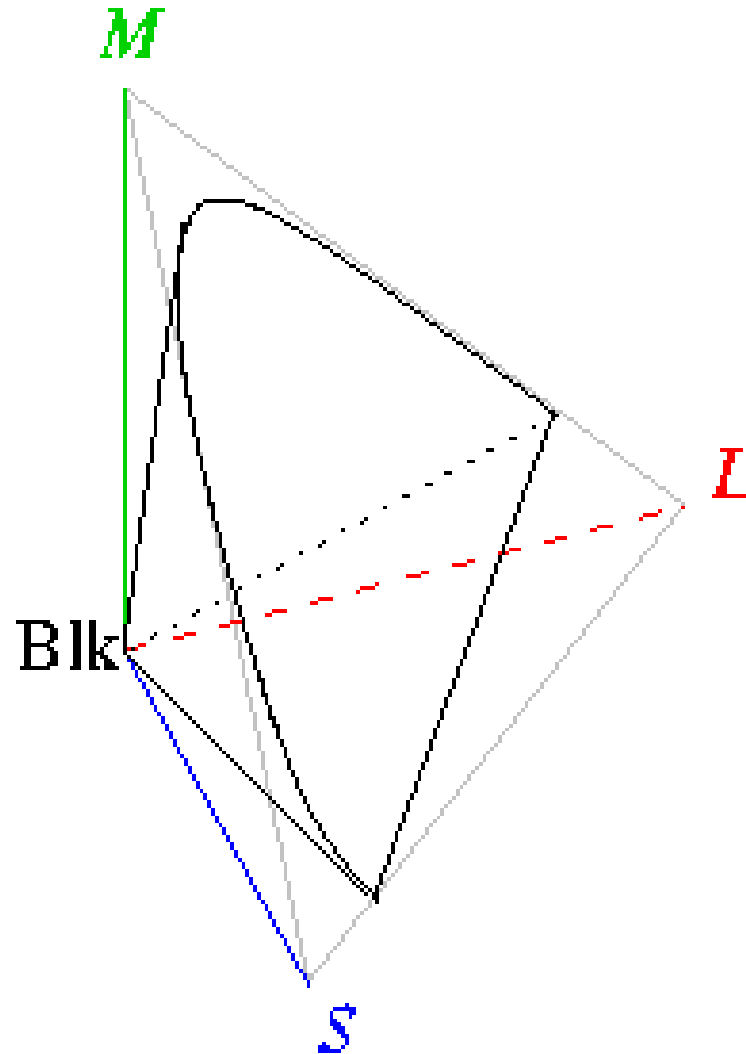
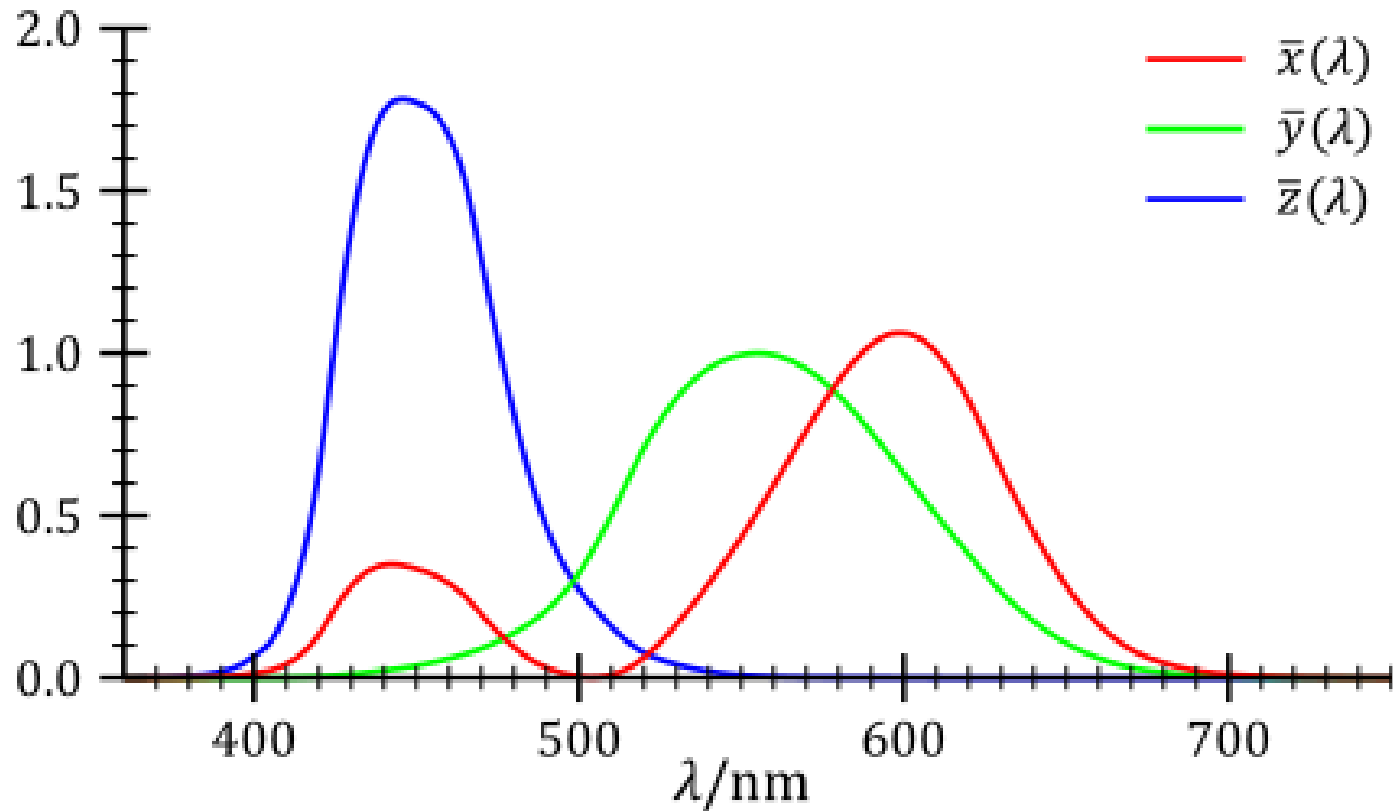


Image by Hankwang, http://en.wikipedia.org/wiki/File:Gamut_full.png

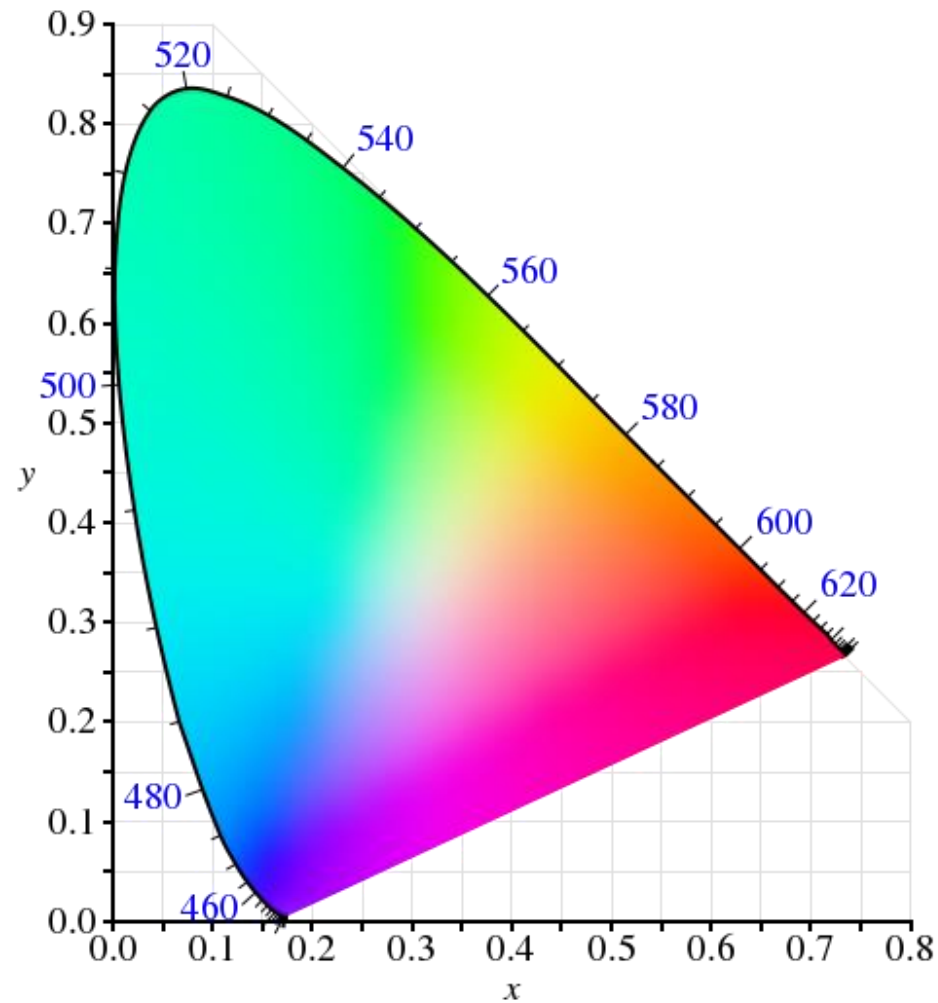
CIE 1931 XYZ



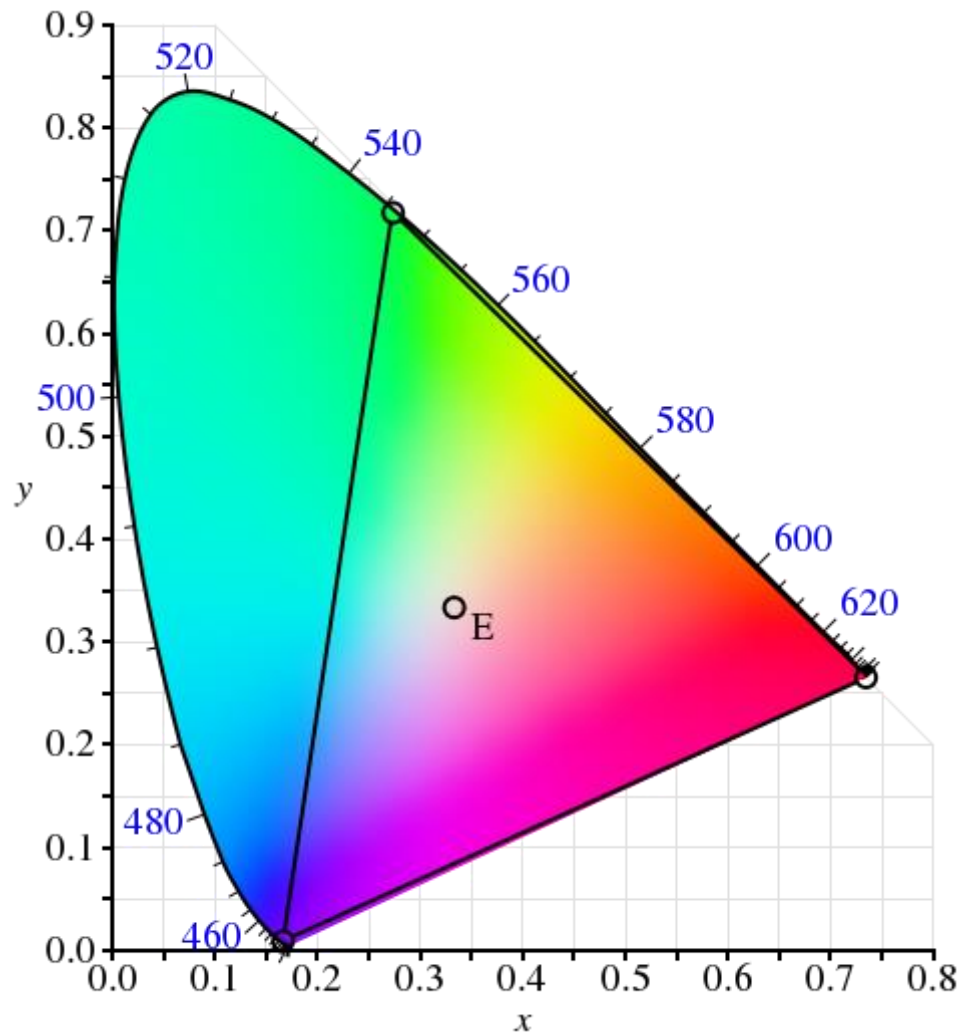
http://en.wikipedia.org/wiki/File:CIE_1931_XYZ_Color_Matching_Functions.svg



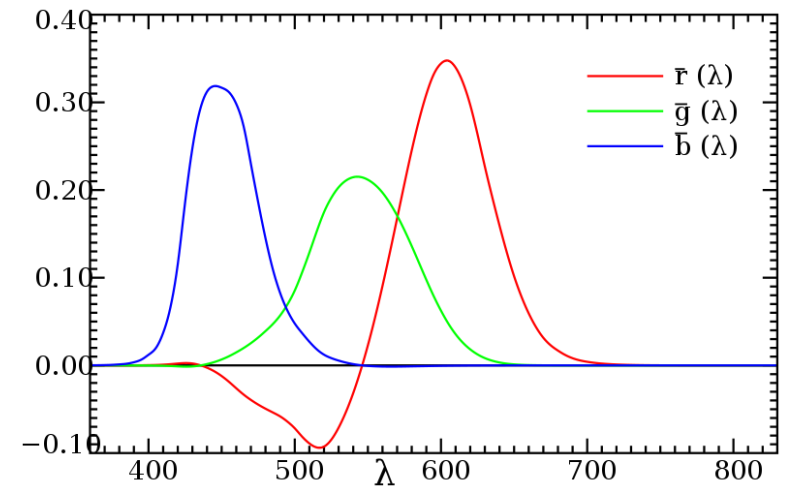
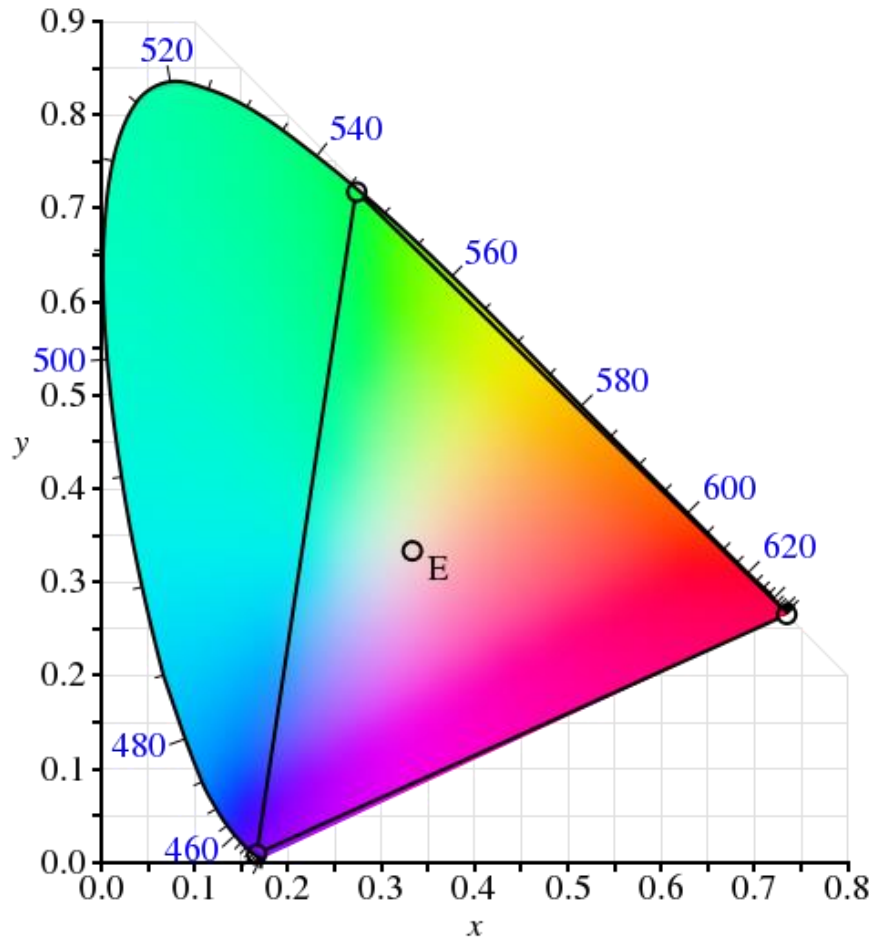
CIE 1931 XYZ, xyY



CIE RGB



CIE RGB

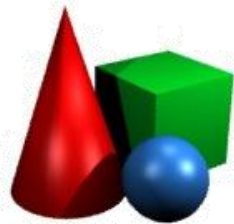


Matching functions



Linearity of color spaces

$$\begin{bmatrix} R \\ G \\ B \end{bmatrix} = \begin{bmatrix} 0.41847 & -0.15866 & -0.082835 \\ -0.091169 & 0.25243 & 0.015708 \\ 0.00092090 & -0.0025498 & 0.17860 \end{bmatrix} \cdot \begin{bmatrix} X \\ Y \\ Z \end{bmatrix},$$



Linearity of color spaces

$$\begin{bmatrix} R \\ G \\ B \end{bmatrix} = \begin{bmatrix} 0.41847 & -0.15866 & -0.082835 \\ -0.091169 & 0.25243 & 0.015708 \\ 0.00092090 & -0.0025498 & 0.17860 \end{bmatrix} \cdot \begin{bmatrix} X \\ Y \\ Z \end{bmatrix},$$

A number of other spaces are obtained as linear transformations of XYZ:

HSV, CMYK, LAB, YUV, ...

In this course we'll only deal with RGB.



Various RGB spaces

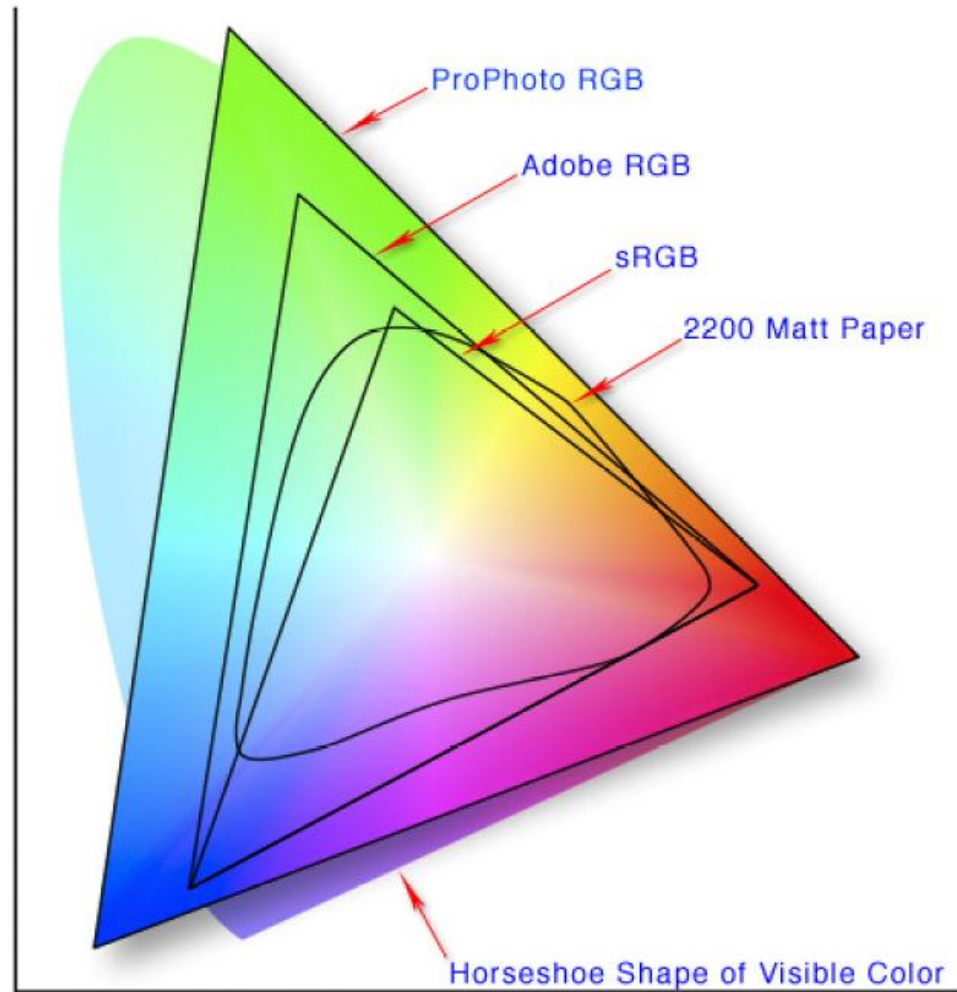


Image by by Jeff Schewe, <http://en.wikipedia.org/wiki/File:Colorspace.png>

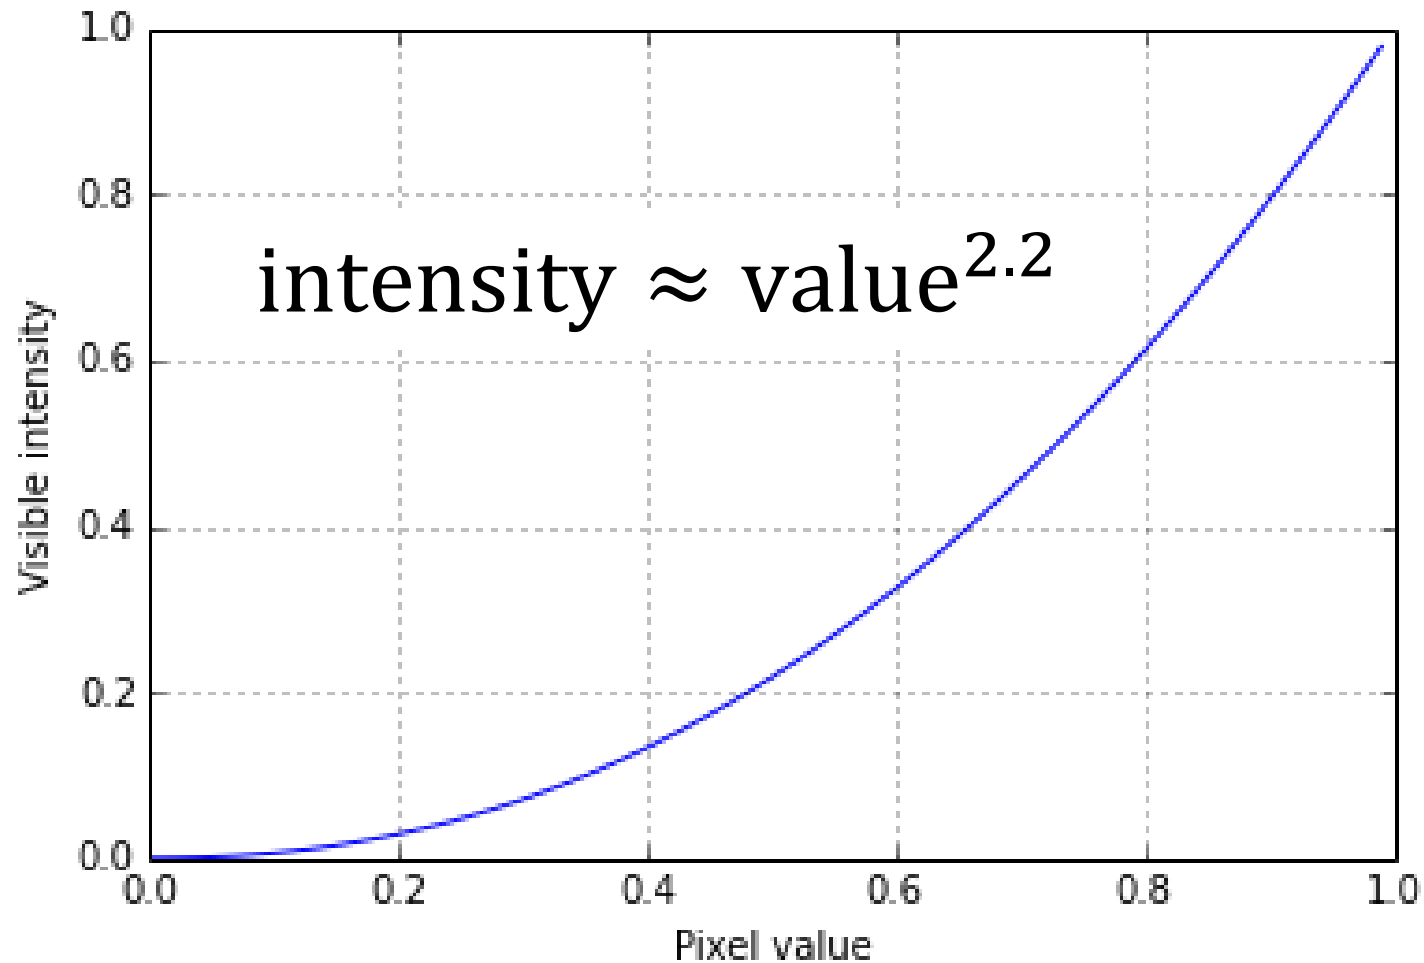


sRGB

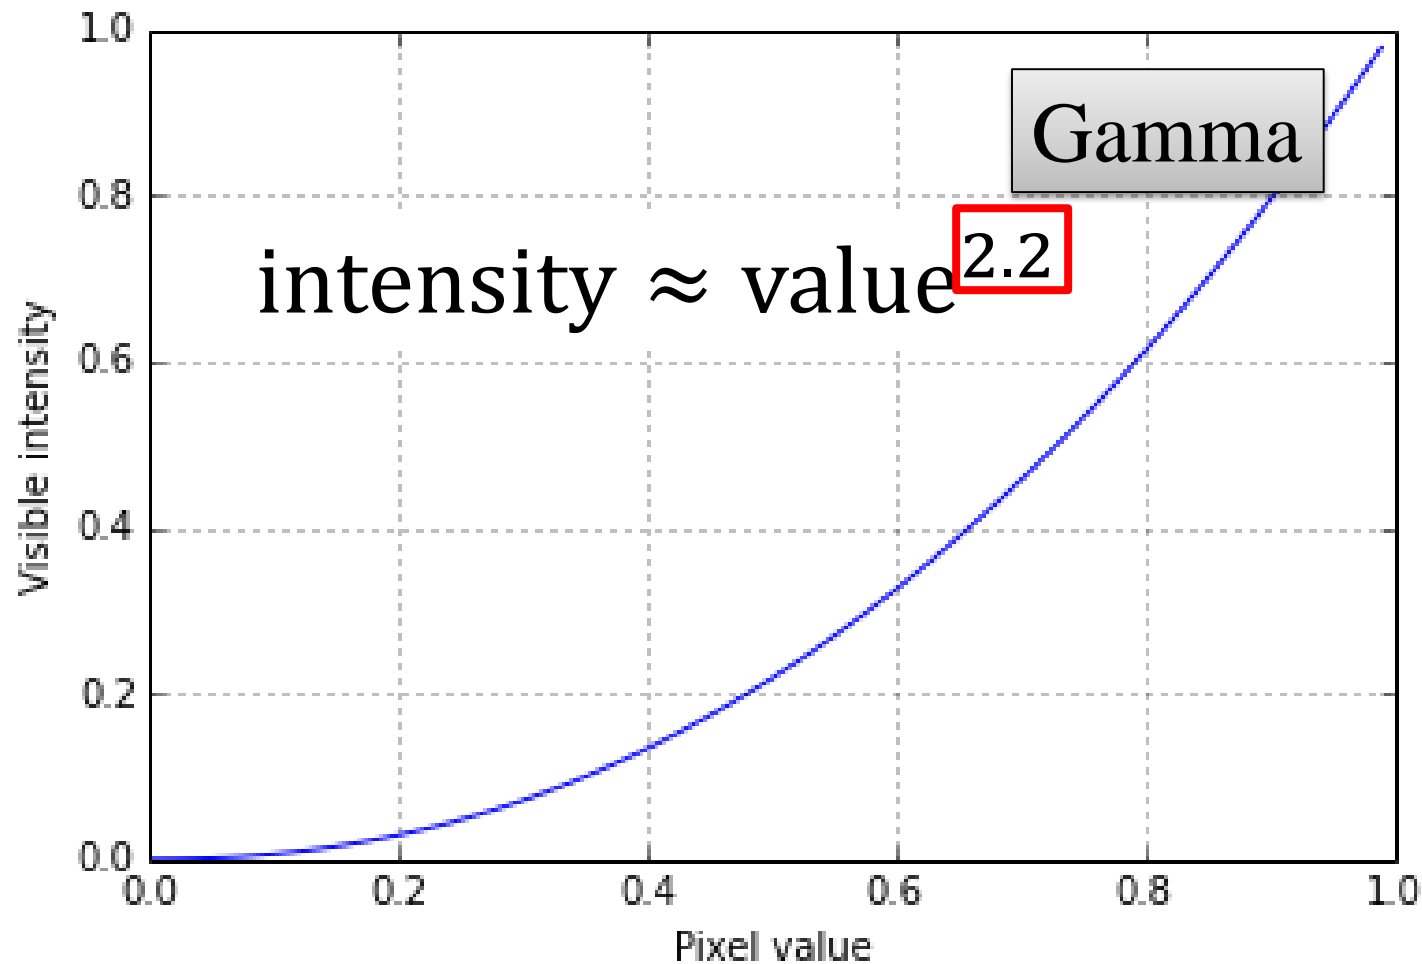
- CIE RGB model is *linear* with respect to light intensities: shining two 0.2 red lights together produces a 0.4 red light.
- However, original CRT monitors were *not* linear in how they work: sending a signal twice as strong would not double light intensity.



CRT response curve



CRT response curve



sRGB

- This led to the standardization of a (non-linear) **sRGB** (“**standard RGB**”) color space, corresponding to the “typical CRT”.
- $\text{CIE RGB} \sim \text{sRGB}^{2.2}$

* Actual transformation slightly more complicated:
see <http://en.wikipedia.org/wiki/SRGB>



sRGB

- This led to the standardization of a (non-linear) **sRGB** (“standard RGB”) color space. **Nearly all modern digital display and imaging devices (cameras, monitors, TVs, printers, scanners, etc) use sRGB.**
- CIE 1931 XYZ color space

* Actual transformation slightly more complicated:
see <http://en.wikipedia.org/wiki/SRGB>



sRGB

- This led to the standardization of a (non-linear) **sRGB** (“standard RGB”) color space, defined by the “sRGB IEC61966-2.1”.

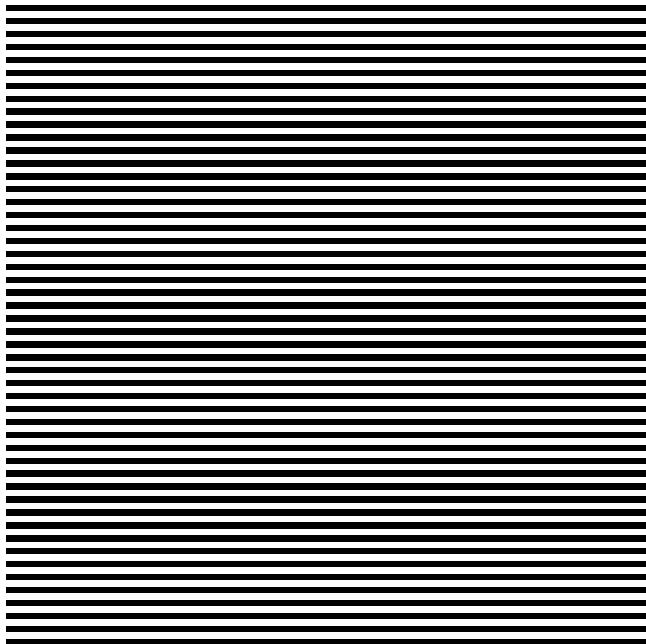
Most images are saved in sRGB

- CIE

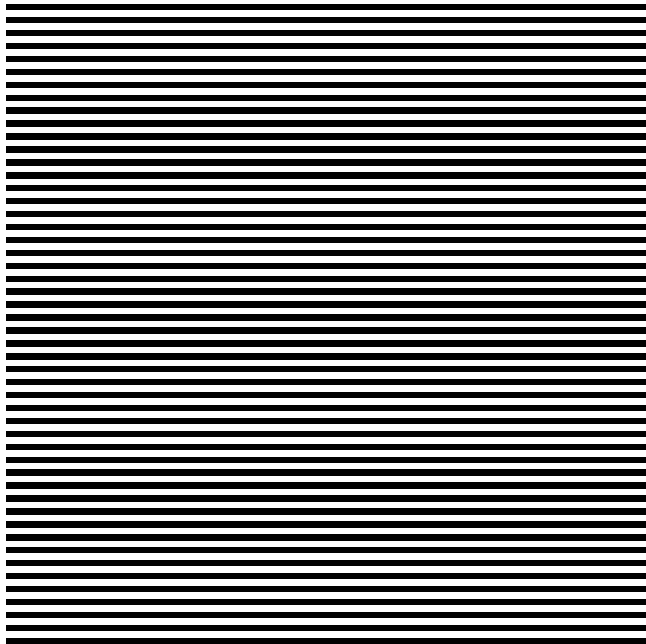
* Actual transformation slightly more complicated:
see <http://en.wikipedia.org/wiki/SRGB>



RGB(0.5) vs. sRGB(0.5)



RGB(0.5) vs. sRGB(0.73)



Gamma correction

- The process of converting from linear intensities to sRGB is called *gamma correction* or *gamma encoding*.
- $\text{sRGB} \sim \text{CIE RGB}^{\frac{1}{2.2}}$
- The inverse operation is called **gamma decoding**



Light modeling

- When rendering an object, we would like to compute the color of a pixel so that it would imitate physical processes of light scattering from the object.



Light modeling

- When rendering an object, we would like to compute the color of a pixel so that it would imitate physical processes of light scattering from the object.
- The actual physics of light is hard to compute. Thus, we shall be using fake approximations instead.



Light modeling



Surface



Light source
with intensity I

Light modeling

Light falls from
a light source

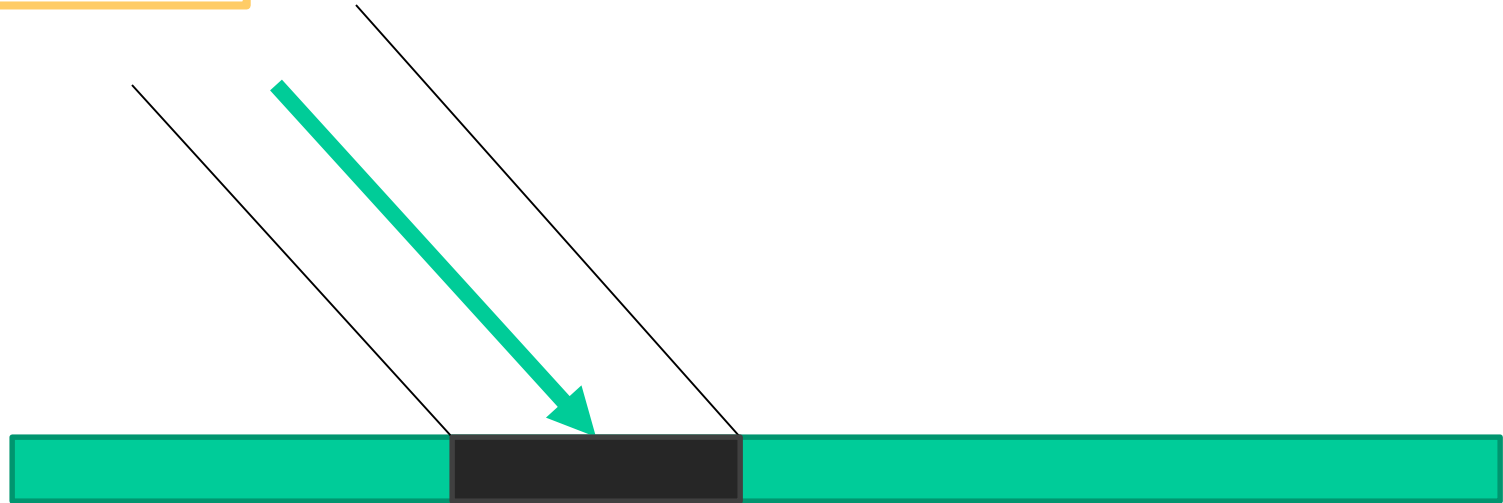


Surface



Light source
with intensity I

Light modeling



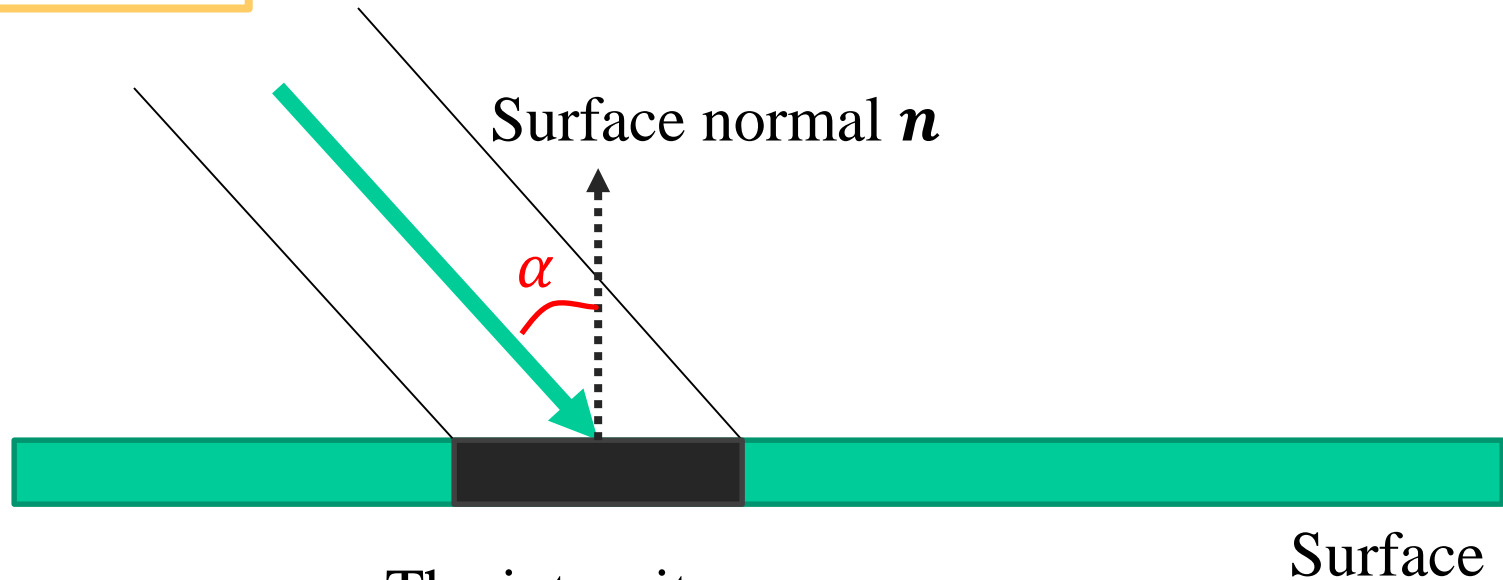
Surface

The more perpendicular
is the light, the greater is its
intensity per area unit.



Light source
with intensity I

Light modeling

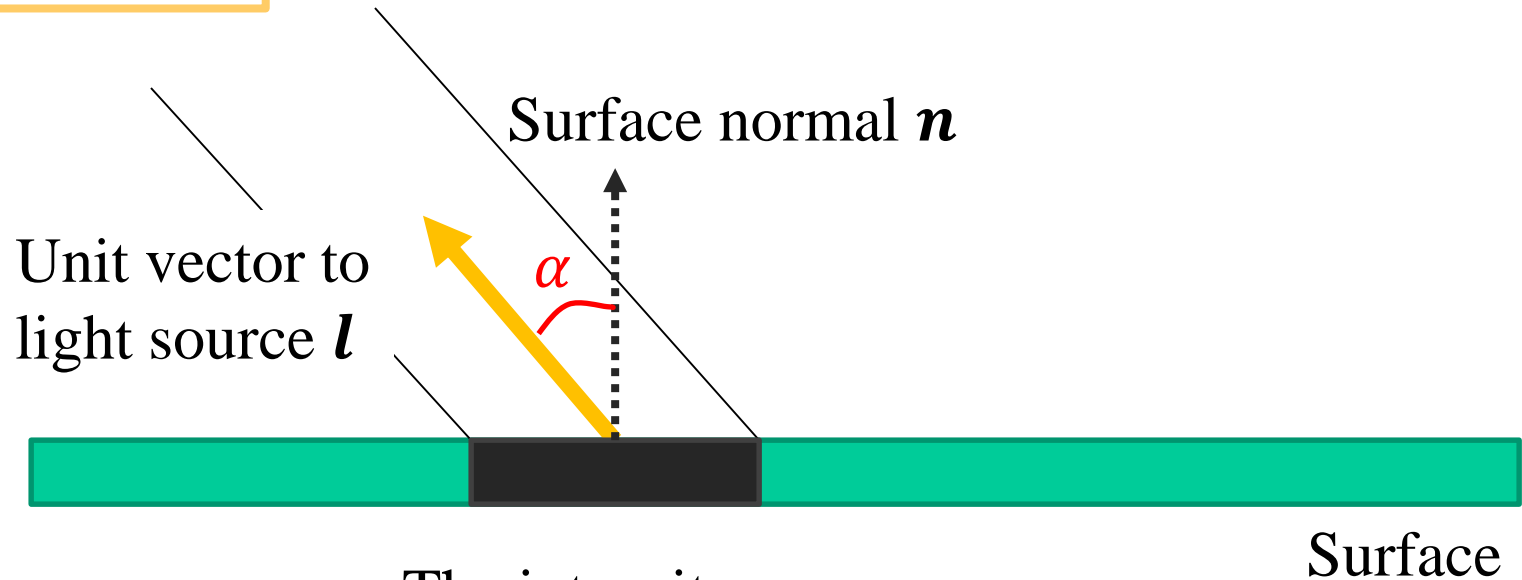


The intensity per
unit area scales
proportionally
to $\cos(\alpha)$



Light source
with intensity I

Light modeling

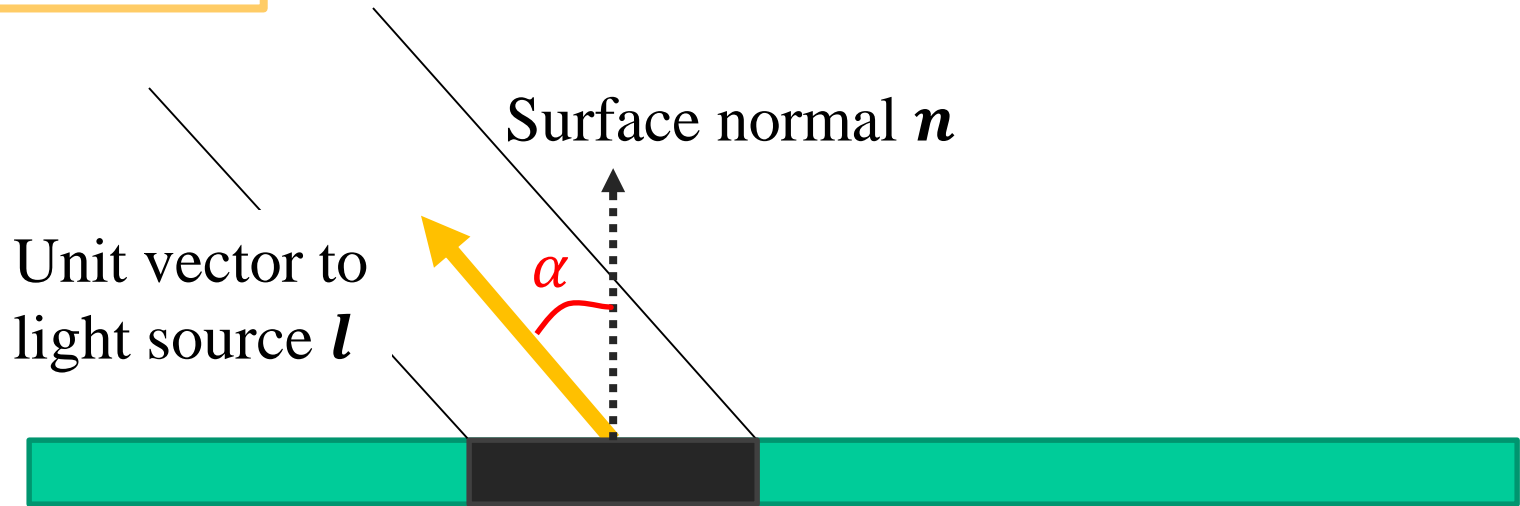


The intensity per
unit area scales
proportionally
to $\cos(\alpha) = l^T n$



Light source
with intensity I

Light modeling



The intensity that reaches
each unit area is thus equal to

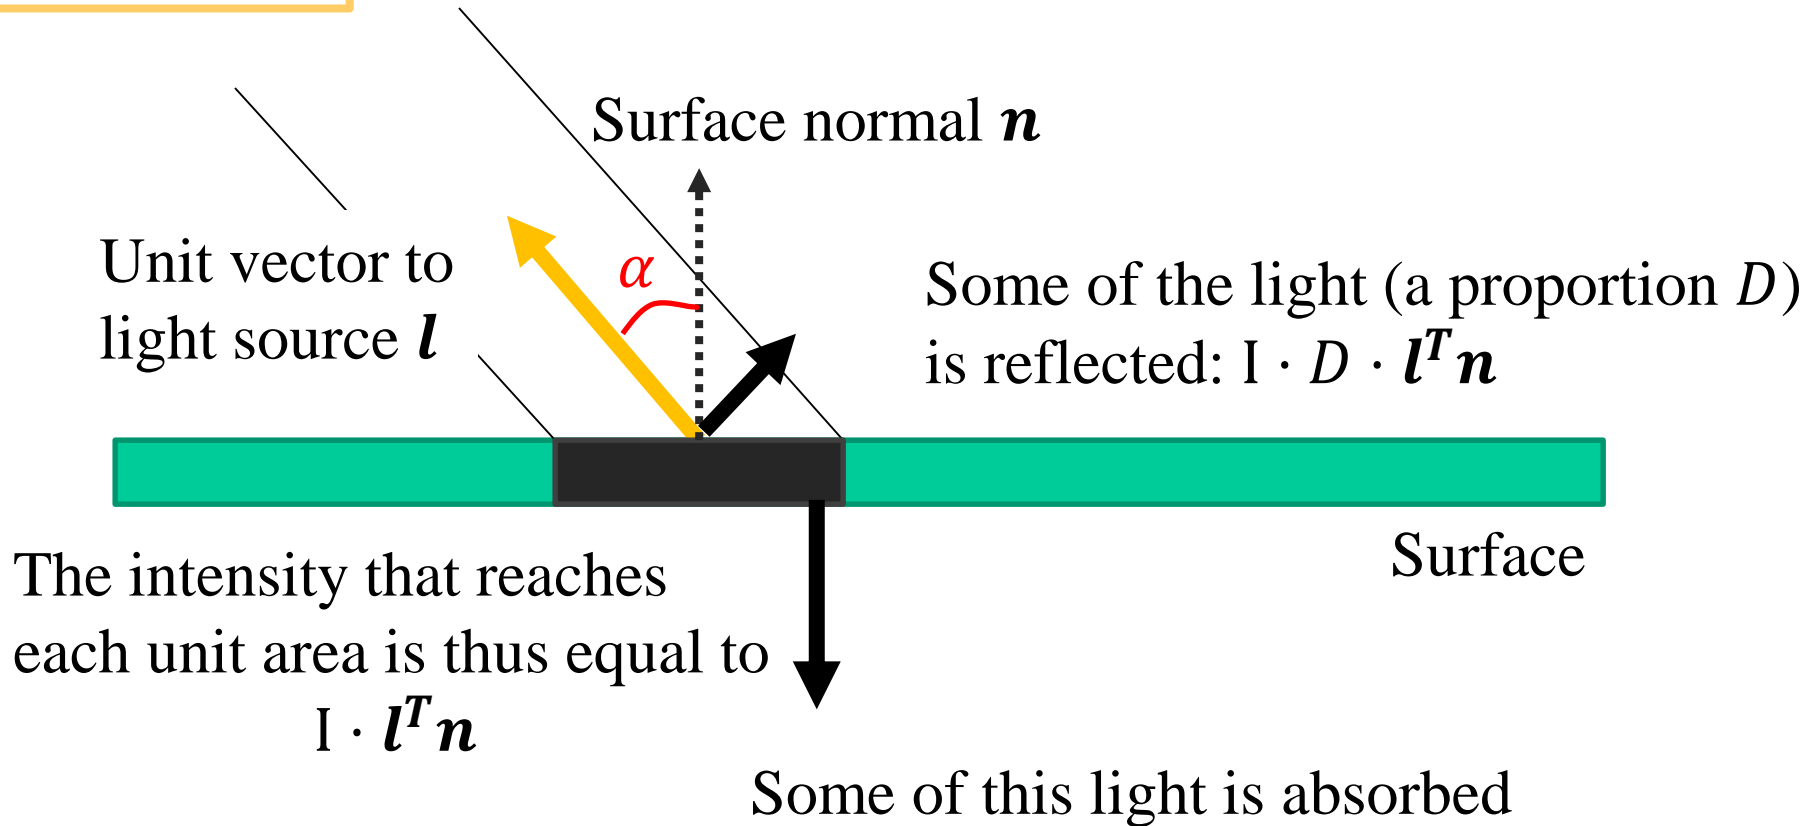
$$I \cdot l^T n$$

Surface



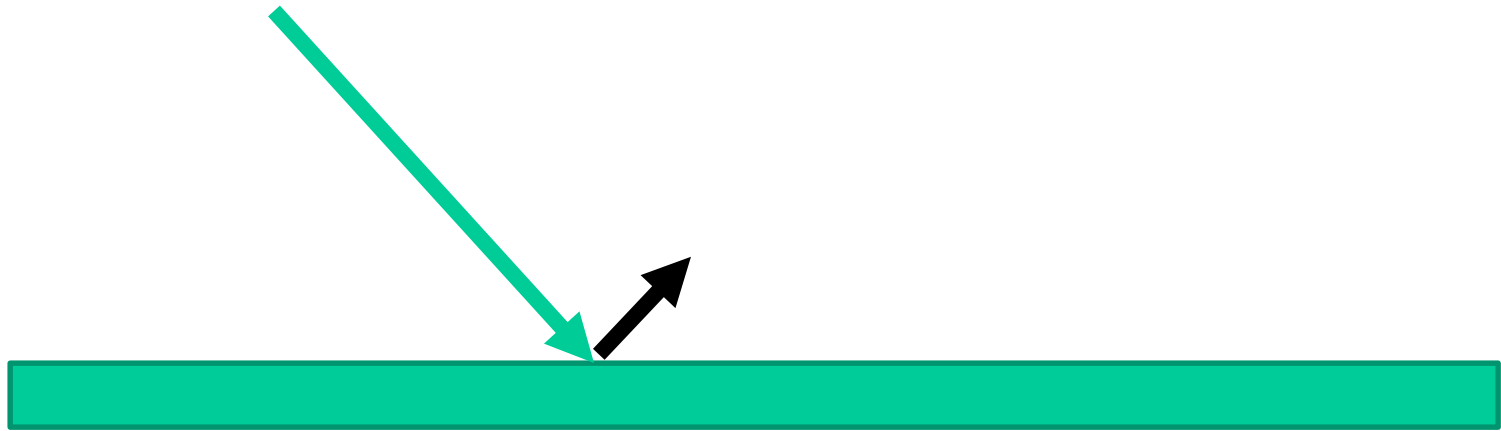
Light source
with intensity I

Light modeling



Light source
with intensity I

Specular reflection

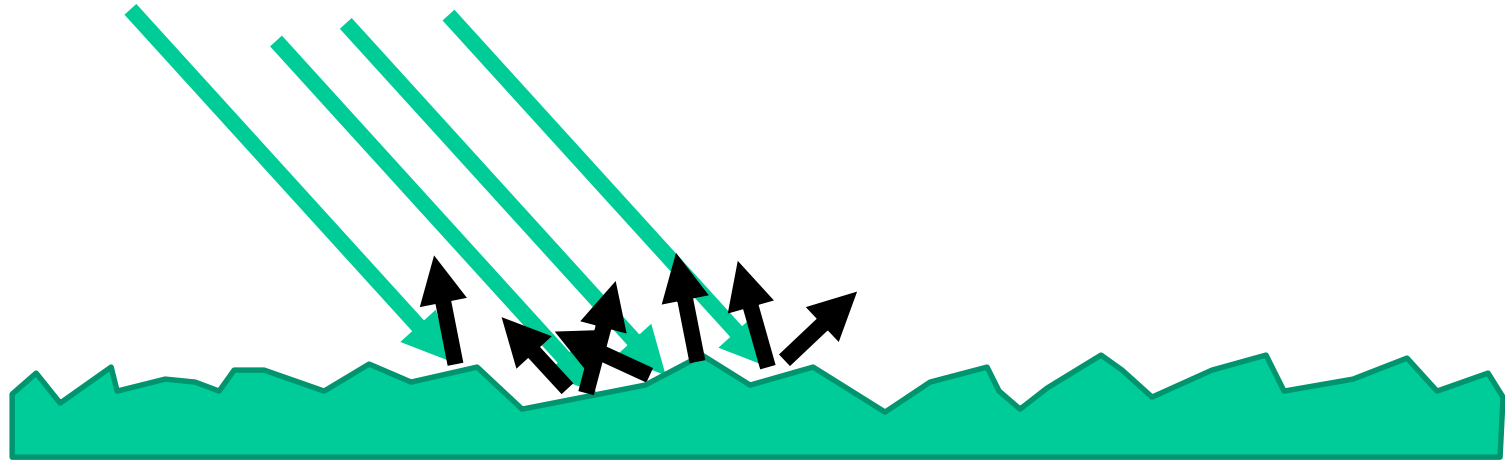


From a perfectly smooth surface,
light is reflected ideally



Light source
with intensity I

Diffuse reflection

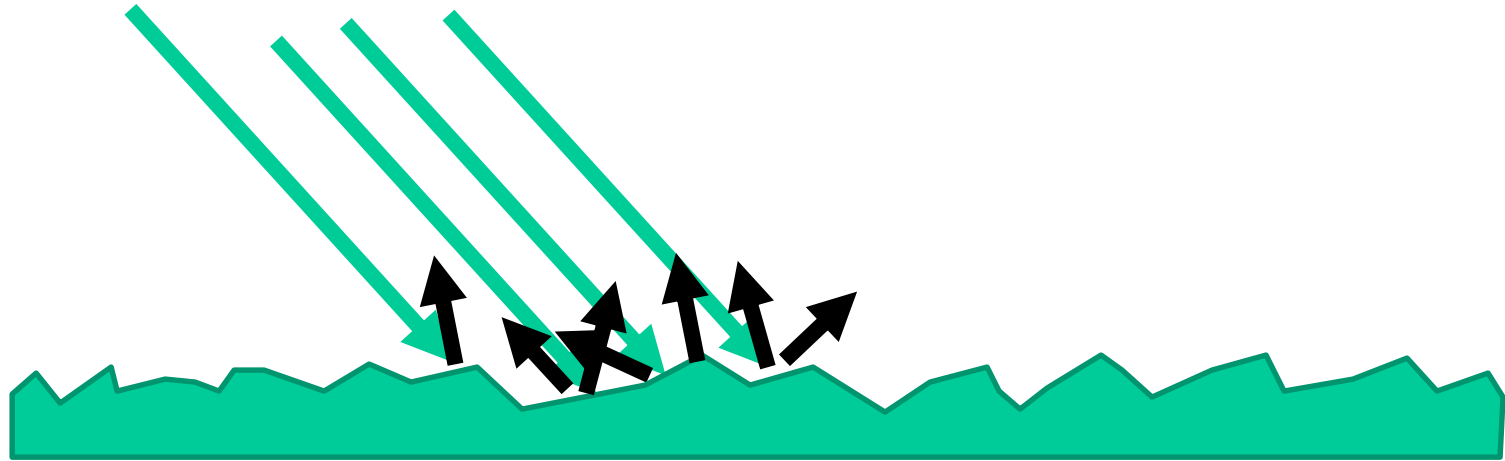


Most surfaces are not perfectly smooth,
and can be regarded as consisting of
microfacets (each of which is a tiny
perfect reflector)



Light source
with intensity I

Diffuse reflection



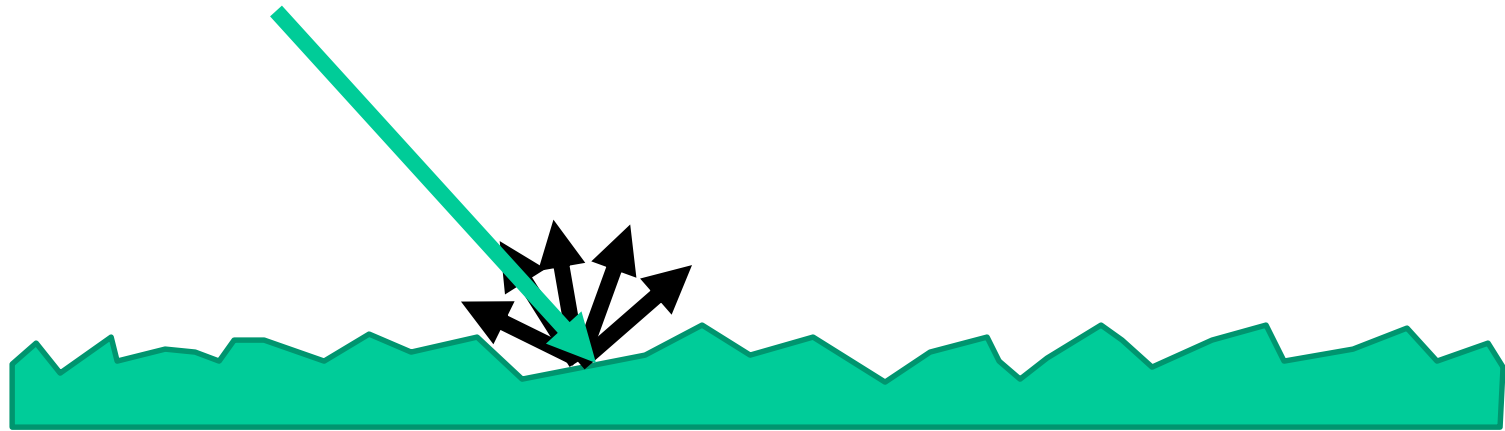
The type of reflection diffusion depends on the distribution of microfacets.

An “ideally” diffusing surface will diffuse light equally in all directions.



Light source
with intensity I

Diffuse reflection



The type of reflection diffusion depends on the distribution of microfacets.

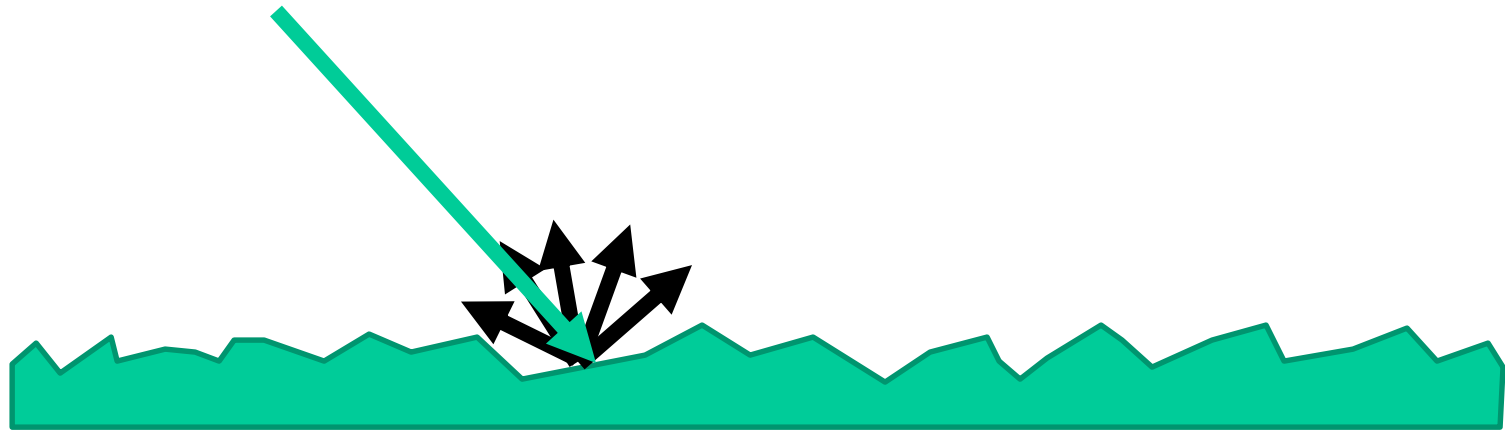
An “ideally” diffusing surface will diffuse light equally in all directions.

* Formally, it is not “equally in all directions”, see http://en.wikipedia.org/wiki/Lambert's_cosine_law



Light source
with intensity I

Lambertian reflection



Lambertian (diffuse) reflection:

The amount of light reaching the observer is

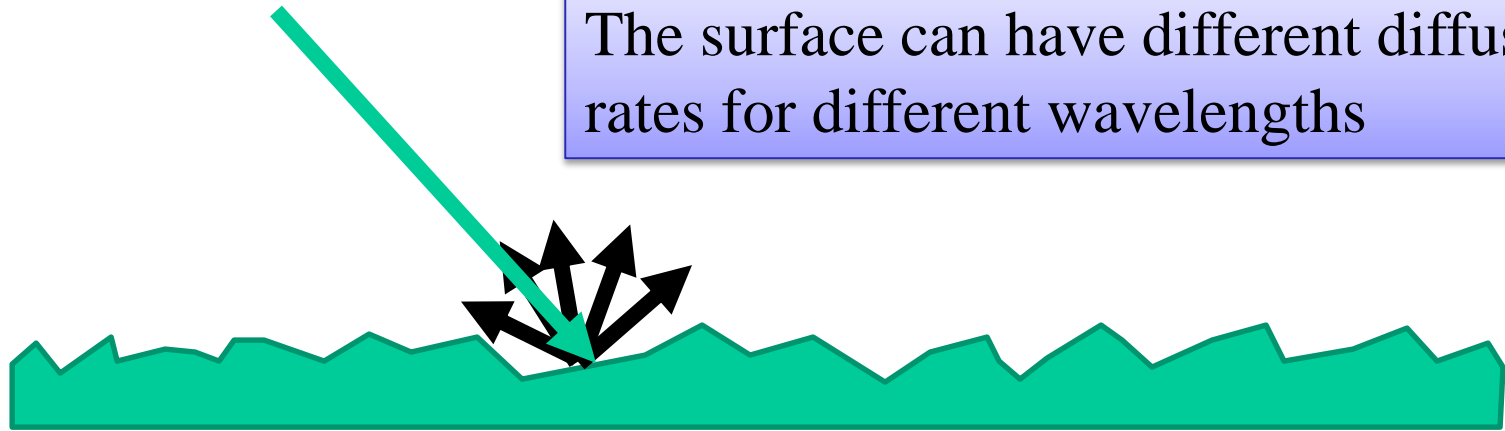
$$I \cdot D \cdot (l^T n)$$



Light source
with intensity I

Lambertian reflection

The surface can have different diffusion rates for different wavelengths



Lambertian (diffuse) reflection:

The amount of light reaching the observer is

$$I \cdot D \cdot (l^T n)$$



Light source
with intensity I

Lambertian reflection

We specify separate diffusion
coefficient for each color component:
 (D_R, D_G, D_B)



Lambertian (diffuse) reflection:

The amount of light reaching the observer is

$$I \cdot D \cdot (l^T n)$$



Lighting in OpenGL

The old-school way:

```
glEnable(GL_LIGHTING);  
glEnable(GL_LIGHT0);  
  
float direction[] = {0, 0, 1, 0};  
glLightfv(GL_LIGHT0, GL_POSITION, direction);  
  
float intensity[] = {1, 1, 1, 1};  
glLightfv(GL_LIGHT0, GL_DIFFUSE, intensity);
```

* Of course it's deprecated, but it is nice & easy for simpler cases.



Lighting in OpenGL

```
float material[] = {1, 0, 0, 1};  
glMaterialfv(GL_FRONT, GL_DIFFUSE, material);
```

or

```
glEnable(GL_COLOR_MATERIAL);  
glColorMaterial(GL_FRONT, GL_DIFFUSE);  
glColor3f(1, 0, 0);
```



Lighting in OpenGL

```
glShadeModel (GL_SMOOTH) ;
```

or

```
glShadeModel (GL_FLAT) ;
```



Lighting in OpenGL

```
glShadeModel (GL_SMOOTH) ;
```

NB: “Old-style” OpenGL does **not** do gamma correction in lighting computations

```
glShadeModel (GL_FLAT) ;
```

* Unless your graphics card supports ARB_framebuffer_sRGB extension.



Lighting in OpenGL

- The new school-way: using GLSL.

```
smooth out vec4 vertex_color;  
uniform vec3 light_dir;  
uniform vec4 diffuse_l, diffuse_m;
```

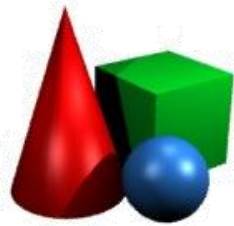
Vertex shader

```
void main(void) {  
    float c = clamp(dot(gl_Normal, light_dir), 0, 1);  
    vertex_color = diffuse_l * diffuse_m * c;  
    gl_Position = ftransform();  
}
```

```
in vec4 vertex_color;
```

```
void main(void) {  
    glFragColor = vertex_color;  
}
```

Fragment shader



Lighting in OpenGL

- The new school-way: using GLSL.

```
smooth out vec4 vertex_color;  
uniform vec3 light_dir;  
uniform vec4 diffuse_l, diffuse_m;
```

Vertex shader

```
void main(void) {  
    float c = clamp(dot(gl_Normal, light_dir), 0.0, 1.0);  
    vertex_color = diffuse_l * diffuse_m * c;  
    vertex_color = pow(vertex_color, 1.0/2.2);  
    gl_Position = transform();  
}
```

Gamma correction

```
in vec4 vertex_color;
```

```
void main(void) {  
    glFragColor = vertex_color;  
}
```

Fragment shader



Per-vertex vs per-fragment lighting

- **Per-vertex lighting:** perform lighting computations in the vertex shader, assign resulting colors to vertices, and simply interpolate them.
- **Per-fragment lighting:** perform lighting computations in the fragment shader.



Per-vertex vs per-fragment lighting

- **Per-vertex lighting:** perform lighting computations in the vertex shader, assign resulting colors to vertices, and simply interpolate across the surface.
NB: “Old-style” OpenGL can only do per-vertex lighting.
- **Per-fragment lighting:** perform lighting computations in the fragment shader.



Ambient light

- To deal with unlit surfaces, we introduce the average level of “ambient” light. It simulates reflections scattered from all over.
- The total light intensity is then:

$$I_A \cdot A + I_D \cdot D \cdot \mathbf{l}^T \mathbf{n}$$



Ambient light in OpenGL

```
glEnable(GL_LIGHTING);  
glEnable(GL_LIGHT0);  
  
glLightfv(GL_LIGHT0, GL_POSITION, direction);  
  
glLightfv(GL_LIGHT0, GL_DIFFUSE, diffuse_l);  
glLightfv(GL_LIGHT0, GL_AMBIENT, ambient_l);  
  
glMaterialfv(GL_FRONT, GL_DIFFUSE, diffuse_m);  
glMaterialfv(GL_FRONT, GL_AMBIENT, ambient_m);
```



Ambient light in OpenGL

```
glEnable(GL_LIGHTING);  
glEnable(GL_LIGHT0);  
  
glLightfv(GL_LIGHT0, GL_POSITION, direction);  
  
glLightfv(GL_LIGHT0, GL_DIFFUSE, diffuse_l);  
// Makes more sense  
glLightModelfv(GL_LIGHT_MODEL_AMBIENT, ambient_l);  
  
glMaterialfv(GL_FRONT, GL_DIFFUSE, diffuse_m);  
glMaterialfv(GL_FRONT, GL_AMBIENT, ambient_m);
```



Ambient light in OpenGL

```
glEnable(GL_LIGHTING);  
glEnable(GL_LIGHT0);  
  
glLightfv(GL_LIGHT0, GL_POSITION, direction);  
  
glLightfv(GL_LIGHT0, GL_DIFFUSE, diffuse_l);  
glLightModelfv(GL_LIGHT_MODEL_AMBIENT, ambient_l);  
  
// Often more convenient  
glColorMaterial(GL_FRONT, GL_AMBIENT_AND_DIFFUSE);  
glColor3f(...);
```



Ambient light in OpenGL

Vertex shader

```
smooth out vec4 vertex_color;
...
uniform vec3 ambient_l, ambient_m;

void main(void) {
    float c = clamp(dot(gl_Normal, light_dir), 0, 1);

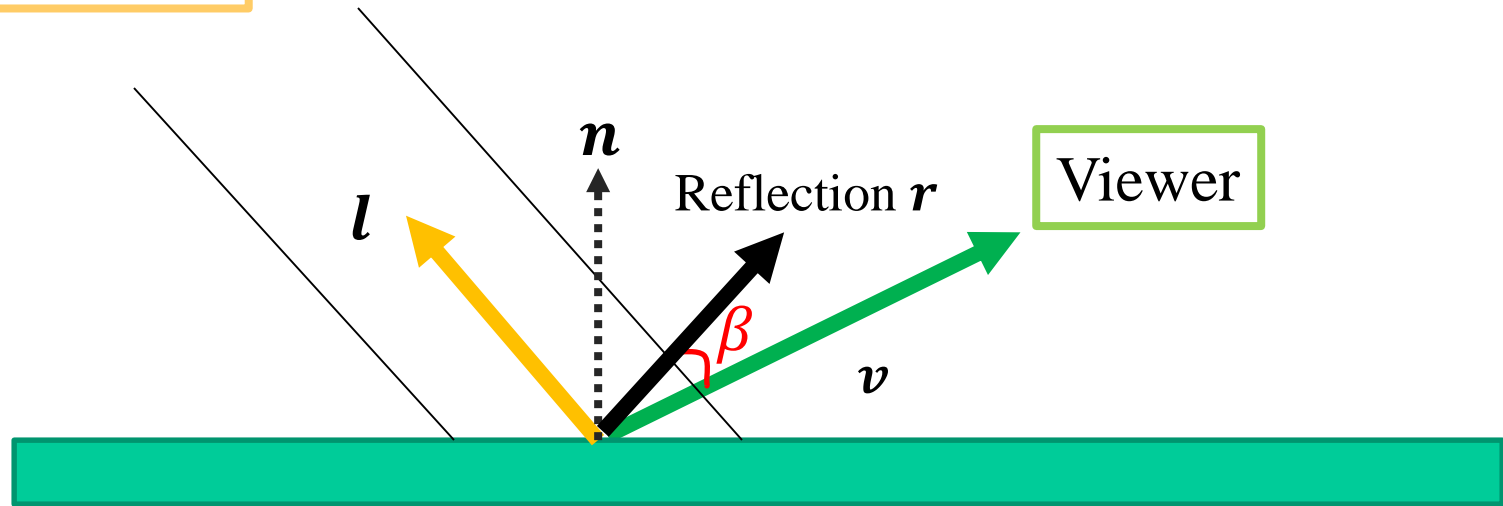
    vertex_color = ambient_l*ambient_m +
                  diffuse_m*diffuse_l*c;
    vertex_color = pow(vertex_color, 1.0/2.2);

    gl_Position = ftransform();
}
```



Light source
with intensity I

Specular reflection

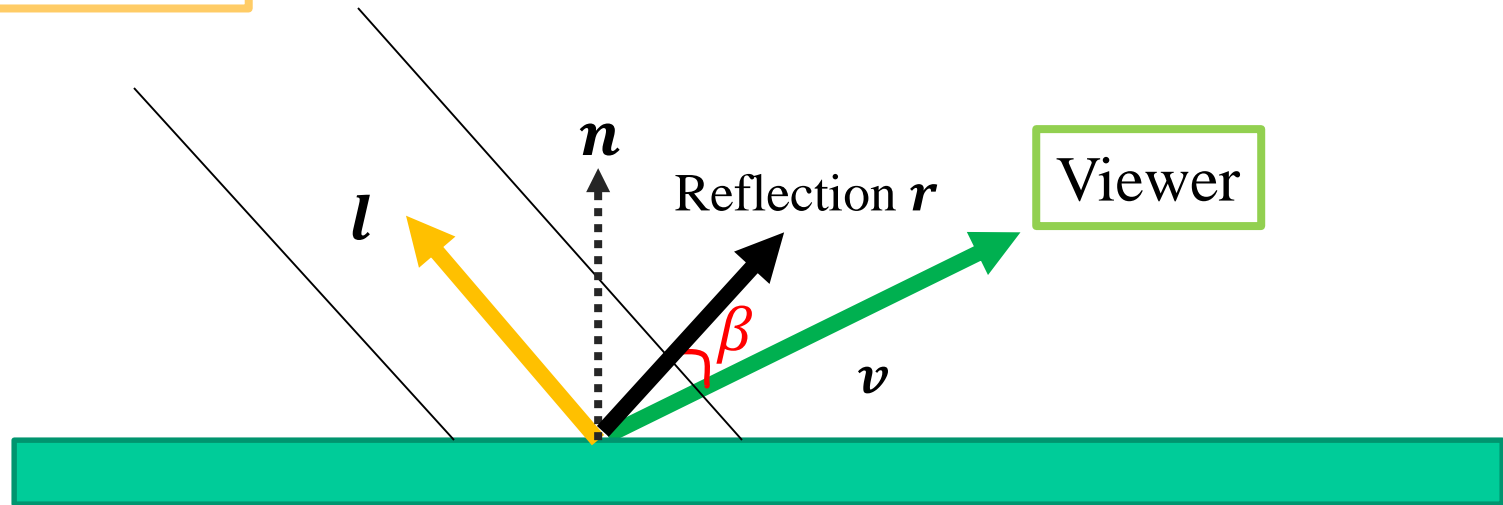


When the viewer is watching in a direction close to that of the reflected light, i.e. when β is small, he should see an intense highlight



Light source
with intensity I

Specular reflection

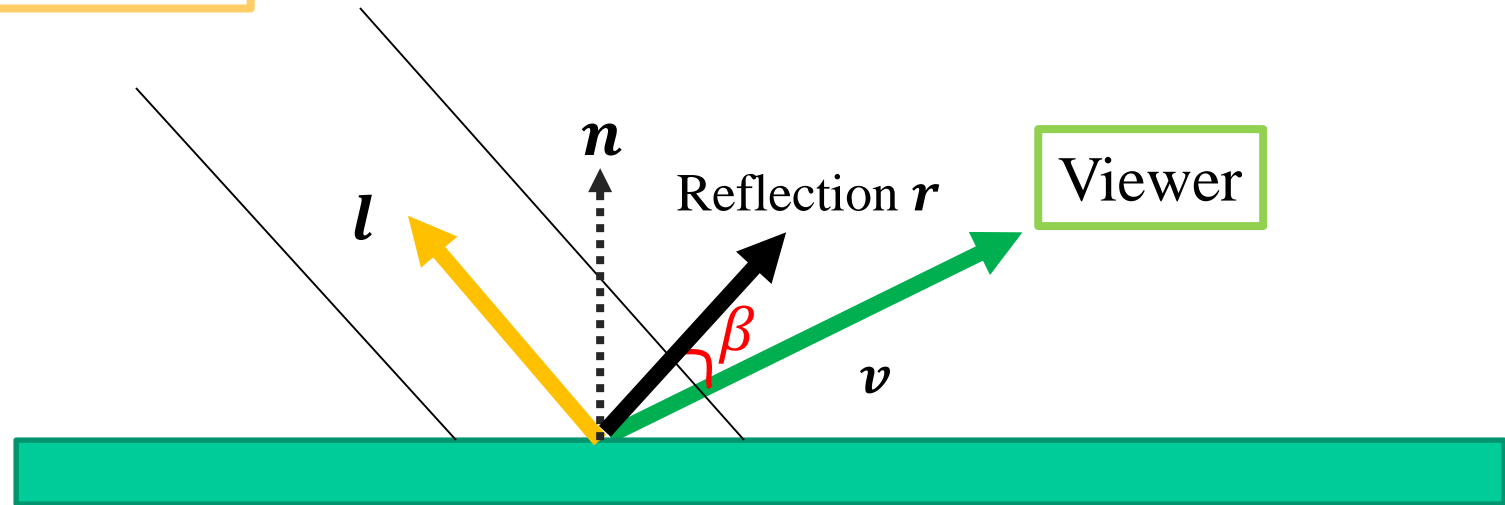


How much the highlight falls off with β depends on the distribution of microfacets.



Light source
with intensity I

Specular reflection



How much the highlight falls off with β depends on the distribution of microfacets.

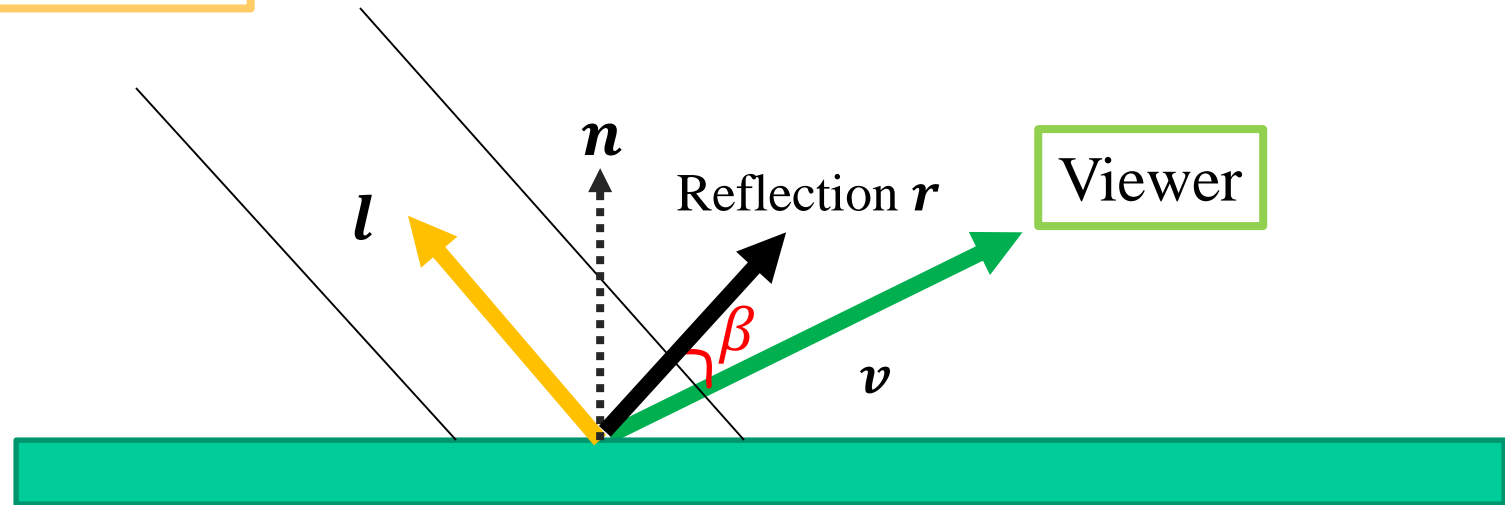
A Gaussian curve is believed to be a good approximation:

$$\text{intensity of specular reflection} \approx e^{-\left(\frac{\beta}{m}\right)^2}$$



Light source
with intensity I

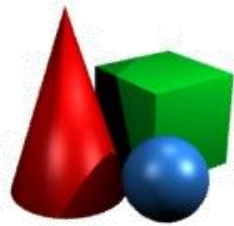
Phong reflection



How much the highlight falls off with β depends on the distribution of microfacets.

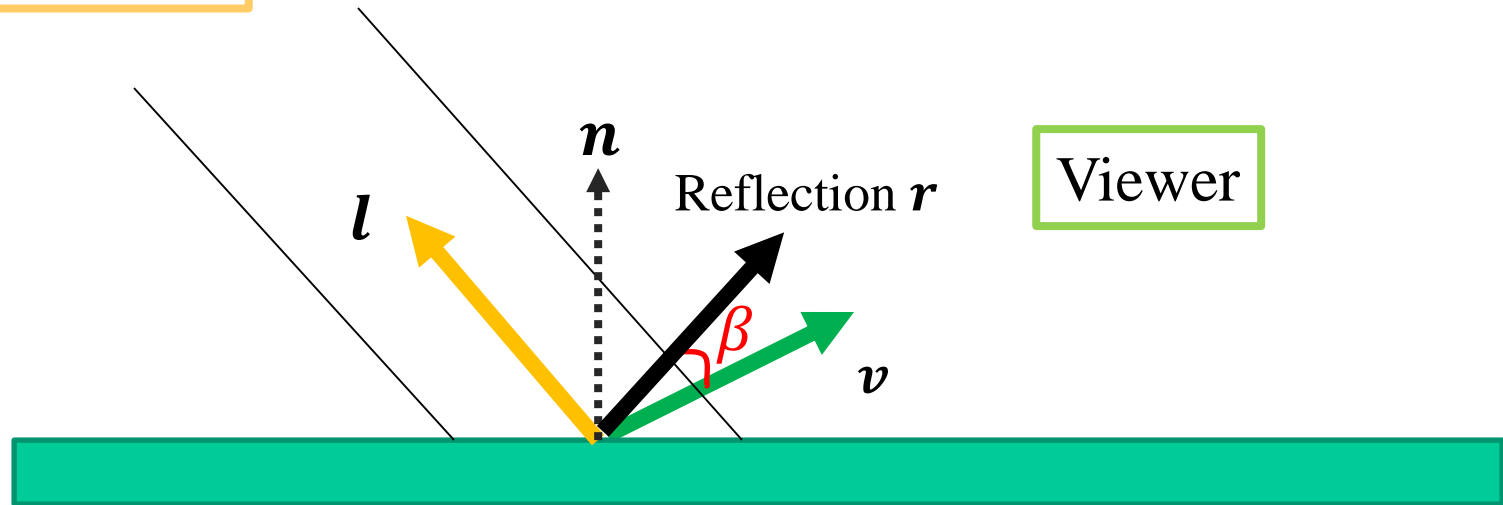
In practice, the following approximation (*Phong model*) is often used:

$$\text{intensity of specular reflection} \approx \cos(\beta)^s$$



Light source
with intensity I

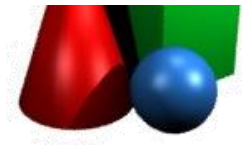
Phong reflection



How much the highlight falls off with β depends on the distribution of microfacets.

In practice, the following approximation (Phong model) is often used:

$$\text{intensity of specular reflection} \approx (r^T v)^s$$



Phong model

- The complete Phong lighting model is thus:

$$I_A \cdot A + I_D \cdot D \cdot \mathbf{l}^T \mathbf{n} + I_S \cdot S \cdot (\mathbf{r}^T \mathbf{v})^s$$



Phong model

- The complete Phong lighting model is thus:

$$I_A \cdot A + I_D \cdot D \cdot \mathbf{l}^T \mathbf{n} + I_S \cdot S \cdot (\mathbf{r}^T \mathbf{v})^s$$

- Note that objects often have different specular and diffuse colors (e.g. red plastic object has white specular highlights).



Lighting in OpenGL

```
glLightfv(GL_LIGHT0, GL_DIFFUSE, diffuse_l);  
glLightfv(GL_LIGHT0, GL_AMBIENT, ambient_l);  
glLightfv(GL_LIGHT0, GL_SPECULAR, specular_l);  
  
glMaterialfv(GL_FRONT, GL_DIFFUSE, diffuse_m);  
glMaterialfv(GL_FRONT, GL_AMBIENT, ambient_m);  
glMaterialfv(GL_FRONT, GL_SPECULAR, specular_m);  
glMaterialf(GL_FRONT, GL_SHININESS, 2.0);
```



Lighting in OpenGL

```
smooth out vec4 position;  
smooth out vec3 normal;
```

Vertex shader

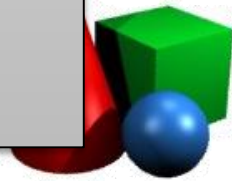
```
void main(void) {  
    position = gl_ModelViewMatrix * gl_Vertex;  
    normal = gl_NormalMatrix * gl_Normal;  
    gl_Position = ftransform();  
}
```



Lighting in OpenGL

Fragment shader

```
in vec4 position;  
in vec3 normal;  
  
uniform vec3 l; // Direction to light  
uniform vec3 v; // Direction to camera  
// Light and material params  
uniform vec3 a_l, a_m, d_l, d_m, s_l, s_m;  
uniform float shininess;  
void main(void) {  
    vec3 n = normalize(normal);  
    float lambert = clamp(dot(n, l), 0, 1);  
    vec3 r = n*2*dot(n, l) - l;  
    float phong = clamp(dot(r, v), 0, 1);  
    phong = pow(phong, shininess);  
  
    gl_FragColor = a_l*a_m + d_l*d_m*lambert  
                  + s_l*s_m*phong;  
}
```



High Dynamic Range correction

- If you add many light components together, you may end up with color values exceeding 1.0.
- By default they are clamped to 1.0.
- In reality human, eye automatically normalizes for maximal brightness.



HDR correction

- For scenes with highly lit objects you should remap colors to fit the 0..1 range.
- The easiest option is simply divide the color by the potentially maximal value – this corresponds to tuning the “aperture” of the camera. It is known as *HDR correction*.



Point light sources

- So far we only considered *directional* light.
- If the light source is a point in space, we must attenuate its intensity based on the *distance* to light.



Point light sources

- Physically correct attenuation is achieved using the inverse square law:

$$\text{intensity at distance } d = \frac{I}{k + d^2}$$



Point light sources

- Sometimes using a more general form of attenuation results in a (physically less correct, but visually nicer) result:

$$I(d) \approx \frac{I}{k_c + k_\ell d + k_q d^2}$$

Where k_c, k_ℓ, k_q are carefully chosen constants.

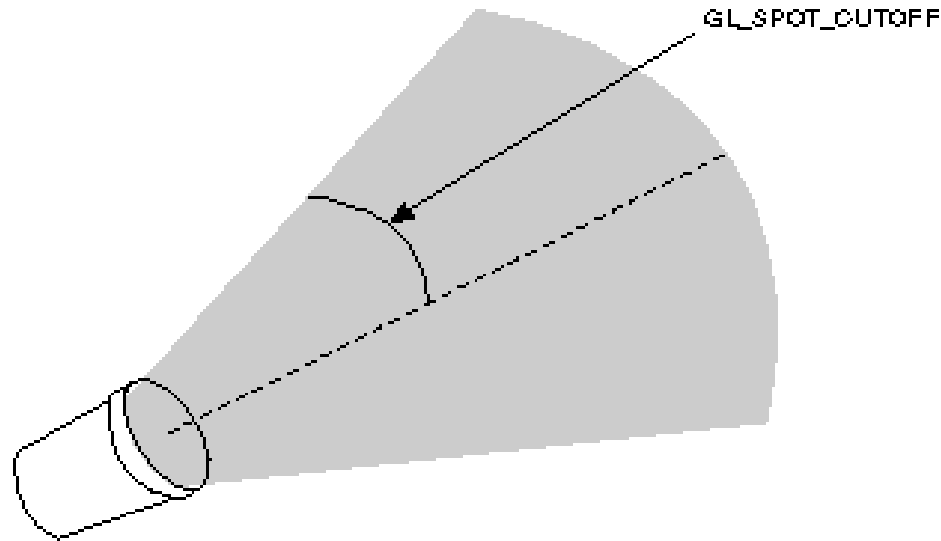


Point light sources

```
float position[] = {-1, 2, 3, 1};  
glLightfv(GL_LIGHT0, GL_POSITION, position);  
  
glLightfv(GL_LIGHT0, GL_CONSTANT_ATTENUATION,  
          1.0);  
glLightfv(GL_LIGHT0, GL_LINEAR_ATTENUATION,  
          1.0);  
glLightfv(GL_LIGHT0, GL_QUADRATIC_ATTENUATION,  
          1.0);
```



Spotlights



```
glLightfv(GL_LIGHT0, GL_SPOT_DIRECTION, dir);  
glLightfv(GL_LIGHT0, GL_SPOT_CUTOFF, 45.0);  
glLightfv(GL_LIGHT0, GL_SPOT_EXPONENT, 1.0);
```



Summary

- Color perception
- CIE XYZ, CIE RGB
- sRGB, Gamma encoding / decoding
- Phong model =
 Ambient, Diffuse, Specular components
- Per-vertex vs Per-fragment lighting
- Fragment shaders
- Directional, point, spotlight light sources
- Gamma correction, HDR correction



Standard Graphics Pipeline

Vertex
transform

Culling and
clipping

Rasterization

Fragment
shading

Visibility tests
& blending

