
Computer Graphics

Textures

Konstantin Tretyakov
kt@ut.ee



Oct 23, 2013

Standard Graphics Pipeline

Vertex
transform

Determine clip-space position of a triangle

Culling and
clipping

Determine whether the triangle is visible

Rasterization

Determine all pixels belonging to the triangle

Fragment
shading

For each pixel, determine its **color**

Visibility tests
& blending

Draw pixel (*if needed*)



Quiz

- Colors are usually described using three numbers, because _____
- Name three different colorspace:

- The “mother” of all colorspace is _____
- What colorspace is used by your digital camera to store pictures?



Quiz

- Colors are usually described using three numbers, because *there are three types of cone cells in our eyes*.
- Name three different colorspace:
e.g. CIE XYZ, CIE RGB, Adobe RGB, CMYK, sRGB, ...
- The “mother” of all colorspace is *CIE XYZ*
- What colorspace is used by your digital camera to store pictures? *sRGB*



Quiz

- Conversion from RGB to sRGB is called
-

and is performed as follows

sRGB := _____



Quiz

- Conversion from RGB to sRGB is called *gamma correction / gamma encoding*

and is performed as follows

$$\text{sRGB} \approx \text{RGB}^{\frac{1}{2.2}}$$



Quiz

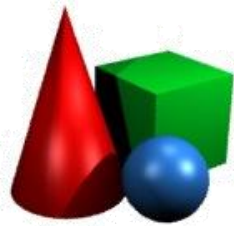
- Phong lighting model
 - Single light component, single light source.
 - Light source intensity (I_a, I_d, I_s)
 - Material parameters (M_a, M_d, M_s), specular shininess 10
 - Quadratic attenuation
-
-



Quiz

- Write down the Phong lighting model for:
 - Single light component, single light source.
 - Light source intensity (I_a, I_d, I_s)
 - Material parameters (M_a, M_d, M_s) , specular shininess 10
 - Quadratic attenuation

$$I_a M_a + \frac{1}{d^2} \left(I_d M_d (\mathbf{n}^T \mathbf{l})_+ + I_s M_s (\mathbf{r}^T \mathbf{v})_+^{10} \right)$$



More basic lighting


- “Blinn specular term”
 - Rather than using $\mathbf{r}^T \mathbf{v}$ in the specular highlight computation, use $\mathbf{n}^T \mathbf{h}$, where \mathbf{h} is the *half-angle*:

$$\mathbf{h} = \text{normalize} \left(\frac{\mathbf{l} + \mathbf{v}}{2} \right)$$



Light source
with intensity I

Phong specular




$$(\cos \beta)^s = (r^T v)^s$$



Light source
with intensity I


Blinn specular



$$(\cos \alpha)^s = (n^T h)^s$$





Gaussian-Blinn specular



$$e^{-\frac{\alpha^2}{\sigma^2}} = e^{-\frac{\arccos(n^T h)^2}{\sigma^2}}$$





More basic lighting

- “Spotlight”, GL_SPOT_XXX
 - Attenuate according to the cosine of the angle between light vector and “spotlight direction”

$$I_a M_a + \frac{1}{d^2} (\dots)$$

becomes

$$I_a M_a + (-l^T \mathbf{spot_dir})_+^{\mathbf{spot_exponent}} \cdot \frac{1}{d^2} (\dots)$$




More basic lighting

- “Emissive material”, GL_EMISSION
 - Simply add a constant “emission intensity” term to the light equation:

$$I_a M_a + \text{attenuation} \cdot (...)$$

becomes

$$\mathbf{M}_e + I_a M_a + \text{attenuation} \cdot (...)$$



More basic lighting

- “Fog”, glEnable(GL_FOG), glFog(..)
 - Weaken color intensity in proportion to distance to the viewer (or fragment depth).

I.e. after computing

$$\text{color} := M_e + I_a M_a + \text{attenuation} \cdot (...)$$

blend the result with the “fog color”:

$$t \cdot \text{color} + (1 - t) \cdot \text{fogColor}$$

Where $t = f(\text{distance})$



Fog functions

- **Linear (GL_LINEAR):**

- $f(d) = \frac{end - d}{end - start}$

- **Exponential (GL_EXP):**

- $f(d) = e^{-\lambda d}$

- **Exponential quadratic (GL_EXP2):**

- $f(d) = e^{-(\lambda d)^2}$



Basic lighting summary

- Phong model
 - Ambient, Diffuse, Specular, Emission
 - Per-vertex & per-fragment
- Half-vector & Blinn specular term
- Gauss specular term
- Fog
 - Linear, Exponential, Exponential quadratic



Standard Graphics Pipeline

Vertex
transform

Determine clip-space position of a triangle

Culling and
clipping

Determine whether the triangle is visible

Rasterization

Determine all pixels belonging to the triangle

Fragment
shading

For each pixel, determine its **color**

Visibility tests
& blending

Draw pixel (*if needed*)



Standard Graphics Pipeline

Vertex
transform

Determine clip-space position of a triangle

Culling and
clipping

Determine whether the triangle is visible

Rasterization

Determine all pixels belonging to the triangle

Fragment
shading

For each pixel, determine its

Fixed color

color

Lighting model

Visibility tests
& blending

Draw pixel (*if needed*)



Standard Graphics Pipeline

Where do we take the colors or lighting model parameters (material, normal, ...) from?

Fixed color

Lighting model

color



Standard Graphics Pipeline

So far:

- 1. Specify globally or per-triangle, or**
- 2. Specify per-vertex and interpolate**

Fixed color

Lighting model

color



Standard Graphics Pipeline

Today:

3. Specify with a much higher granularity using a *texture*.

Fixed color

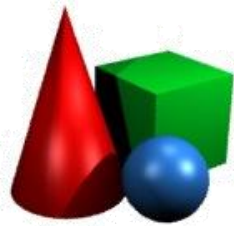

Lighting model

color



Example: Textured triangle

- We can specify color of every pixel of a triangle by mapping it from a *texture image*.



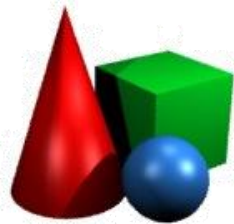
Below the tip of the iceberg

- How attribute interpolation is *actually* done
- Texture filtering
 - Bilinear, mipmapping, anisotropic
- Textures & OpenGL
- Textures beyond images:
 - Precomputation & look-up tables
 - Normal maps, reflection maps, shadow maps
- Procedural textures




Below the tip of the iceberg

- How attribute interpolation is *actually* done
- Texture filtering
 - Bilinear, mipmapping, anisotropic
- Textures & OpenGL
- Textures beyond images:
 - Precomputation & look-up tables
 - Normal maps, reflection maps, shadow maps
- Procedural textures




Linear interpolation


$$A \quad tA + (1 - t)B \quad B$$





Linear interpolation




Linear interpolation




Linear interpolation



Linear interpolation




Linear interpolation




GLSL: `mix(10, 20, 0.3)`




Linear interpolation




Linear interpolation




Linear interpolation




Linear interpolation



Linear interpolation




Linear interpolation



Linear interpolation


- Linear & affine transformations respect linear interpolation:

$$f(t\mathbf{A} + (1 - t)\mathbf{B}) = tf(\mathbf{A}) + (1 - t)f(\mathbf{B})$$




Linear interpolation

- Perspective projection does **not** respect linear interpolation:





Linear interpolation



- Consequently, interpolating in window or clip-space **is not the same** as in world-space.



Linear interpolation



Window-space interpolation



Perspective correct interpolation


Let points A, B be perspectively projected to window-coordinates A^*, B^* . Let

$$C^* = tA^* + (1 - t)B^*$$

Let attribute values F_A and F_B be associated with A and B , varying linearly inbetween.

Then for the corresponding point C :

$$\frac{F_C}{z_C} = t \frac{F_A}{z_A} + (1 - t) \frac{F_B}{z_B}$$



Perspective correct interpolation

Then

$$\frac{F_C}{Z_C} = t \frac{F_A}{Z_A} + (1 - t) \frac{F_B}{Z_B}$$

In particular,

$$\frac{1}{Z_C} = t \frac{1}{Z_A} + (1 - t) \frac{1}{Z_B}$$



Perspective correct interpolation

Thus:

$$F_C = \frac{t \frac{F_A}{Z_A} + (1 - t) \frac{F_B}{Z_B}}{t \frac{1}{Z_A} + (1 - t) \frac{1}{Z_B}}$$



Perspective correct interpolation

If we keep the homogeneous w term before the perspective division, we can do:

$$F_C = \frac{t \frac{F_A}{w_A} + (1 - t) \frac{F_B}{w_B}}{t \frac{1}{w_A} + (1 - t) \frac{1}{w_B}}$$

(this makes interpolation logic independent of what projection matrix you set up).



Interpolation in GLSL

- Vertex shaders output attributes, that are interpolated over the triangle and sent to the fragment shader.
- You can specify the kind of interpolation:

```
flat          out vec3 not_interpolated;  
smooth       out vec3 perspective_correct;  
noperspective out vec3 affine_interpolation;
```

(in the first case, just one of the three vertex attribute values is chosen for all fragments of the triangle)



Texturing

- How attribute interpolation is *actually* done
- Texture filtering
 - Bilinear, mipmapping, anisotropic
- Textures & OpenGL
- Textures beyond images:
 - Precomputation & look-up tables
 - Normal maps, reflection maps, shadow maps
- Procedural textures




Texturing

- How attribute interpolation is *actually* done
- Texture filtering
 - Bilinear, mipmapping, anisotropic
- Textures & OpenGL
- Textures beyond images:
 - Precomputation & look-up tables
 - Normal maps, reflection maps, shadow maps
- Procedural textures




Texture filtering

Consider how a pixel on the screen maps to a texture image



Texture filtering

What texel (texture pixel) should we use for this screen pixel?




Screen pixels

Texture pixels




Texture filtering

1. Pick the nearest one




Texture filtering

2. Bilinearly interpolate (i.e. average) four surrounding pixels.




Texture filtering

Consider how four nearby pixels map to the texture:





Texture filtering

Now we would like to average over a much larger piece of texture



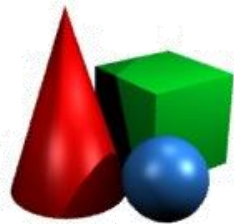

Mipmaps

- To average over large regions of the texture, we pre-average our original texture by computing a set of downscaled versions:





Mipmaps

Level 1 mipmap is downscaled 2x,
so its each pixel is an average of 4
original pixels





Mipmaps

Level 2 mipmap is downscaled 4x,
so its each pixel is an average of 16
original pixels




Mipmaps

Level 3 mipmap is downscaled 8x,
so its each pixel is an average of 64
original pixels




Mipmaps

... etc all the way down to a single-pixel image, which is the overall average of the original image




Texture filtering

Using mipmaps we can efficiently average, e.g. the highlighted piece could be looked up in a level-2 mipmap:




Mipmap-based filtering

- Figure out what region the screen pixel must cover in the texture
- Look up the appropriate average from the mipmaps



Mipmap-based filtering

Option 1: Find the scale-wise closest mipmap, and pick a single pixel from it.




Pixel from a level-2 mipmap covers this average.



Mipmap-based filtering


Option 2: Find the scale-wise closest mipmap, pick several pixels from it and interpolate bilinearly as if it was original image



Mipmap-based filtering


Option 3: Use two different mipmap levels

We can use results from level 1 and level 2-mipmaps, giving perhaps more weight to level 1




Mipmap-based filtering

Option 4: Be very smart and use many mipmap levels to compute the best approximation to the required average:



Mipmap-based filtering

This approach allows averaging over non-square areas, and is thus referred to as *anisotropic filtering*.



Texture filtering

- `GL_NEAREST`
- `GL_LINEAR`
- `GL_LINEAR_MIPMAP_NEAREST`
- `GL_LINEAR_MIPMAP_LINEAR`
- `GL_TEXTURE_MAX_ANISOTROPY_EXT = (number)`



Texturing

- How attribute interpolation is *actually* done
- Texture filtering
 - Bilinear, mipmapping, anisotropic
- Textures & OpenGL
- Textures beyond images:
 - Precomputation & look-up tables
 - Normal maps, reflection maps, shadow maps
- Procedural textures



Texturing

- How attribute interpolation is *actually* done
- Texture filtering
 - Bilinear, mipmapping, anisotropic
- Textures & OpenGL
- Textures beyond images:
 - Precomputation & look-up tables
 - Normal maps, reflection maps, shadow maps
- Procedural textures



Textures & OpenGL

Loading texture image
and mipmaps

```
GLuint textureHandle;  
glGenTextures(1, &textureHandle);  
glBindTexture(GL_TEXTURE_2D, textureHandle);  
  
glTexImage2D(GL_TEXTURE_2D, 0, GL_SRGB8, w, h,  
             0, GL_RGB, GL_UNSIGNED_BYTE, image);  
glTexImage2D(GL_TEXTURE_2D, 1, GL_SRGB8, w/2,h/2,  
             0, GL_RGB, GL_UNSIGNED_BYTE, mipmap1);  
glTexImage2D(GL_TEXTURE_2D, 2, GL_SRGB8, w/4,h/4,  
             0, GL_RGB, GL_UNSIGNED_BYTE, mipmap2);  
  
...
```



Textures & OpenGL

Loading texture image
and mipmaps

```
GLuint textureHandle;  
glGenTextures(1, &textureHandle);  
glBindTexture(GL_TEXTURE_2D, textureHandle);  
  
gluBuild2DMipmaps(GL_TEXTURE_2D, GL_SRGB8, w, h,  
                  GL_RGB, GL_UNSIGNED_BYTE, image);
```



Textures & OpenGL

Specifying filtering method

```
glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_MIN_FILTER,  
                GL_LINEAR_MIPMAP_LINEAR);  
glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_MAG_FILTER,  
                GL_LINEAR);  
glTexParameterf(GL_TEXTURE_2D,  
                GL_TEXTURE_MAX_ANISOTROPY_EXT, 16.0f);
```



Textures & OpenGL

Using the texture


```
glEnable(GL_TEXTURE2D);  
glBindTexture(GL_TEXTURE2D, textureHandle);  
  
glBegin(GL_TRIANGLES);  
    glTexCoord2f(0.0, 0.0);  
    glVertex3f(0.0, 0.0, 0.0);  
    glTexCoord2f(0.0, 1.0);  
    glVertex3f(1.0, 0.0, 0.0);  
    glTexCoord2f(1.0, 1.0);  
    glVertex3f(1.0, 1.0, 0.0);  
glEnd();
```



Textures & OpenGL

Automatically generating coordinates

```
glEnable(GL_TEXTURE_GEN_S);  
glEnable(GL_TEXTURE_GEN_T);  
  
glTexGeni(GL_S, GL_TEXTURE_GEN_MODE, GL_OBJECT_LINEAR);  
glTexGeni(GL_T, GL_TEXTURE_GEN_MODE, GL_OBJECT_LINEAR);  
  
// or  
  
glTexGeni(GL_S, GL_TEXTURE_GEN_MODE, GL_EYE_LINEAR);  
glTexGeni(GL_S, GL_TEXTURE_GEN_MODE, GL_EYE_LINEAR);
```



Textures & OpenGL

How to interpret texture coordinates beyond $[0, 1]$?

```
glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_WRAP_S,  
                GL_REPEAT);  
glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_WRAP_T,  
                GL_REPEAT);
```




Textures & OpenGL

How to combine texture with lighting?

```
glTexEnvi(GL_TEXTURE_ENV, GL_TEXTURE_ENV_MODE,  
          GL_MODULATE);
```

```
glTexEnvi(GL_TEXTURE_ENV, GL_TEXTURE_ENV_MODE,  
          GL_REPLACE);
```

```
glTexEnvi(GL_TEXTURE_ENV, GL_TEXTURE_ENV_MODE,  
          GL_ADD);
```



Textures & GLSL

Texture in the shader is accessed via a “sampler” object


```
uniform sampler2D myTexture;  
  
void main() {  
    gl_FragColor = texture2D(myTexture, gl_TexCoord[0]);  
}
```



Textures & GLSL

```
uniform sampler2D myTexture;  
  
void main() {  
    gl_FragColor = texture2D(myTexture, gl_TexCoord[0]);  
}
```

Of course, it is possible to define multiple sampler objects, each with its own parameters.



Texturing

- How attribute interpolation is *actually* done
- Texture filtering
 - Bilinear, mipmapping, anisotropic
- Textures & OpenGL
- Textures beyond images:
 - Precomputation & look-up tables
 - Normal maps, reflection maps, shadow maps
- Procedural textures



Next week

- How attribute interpolation is *actually* done
- Texture filtering
 - Bilinear, mipmapping, anisotropic
- Textures & OpenGL
- Textures beyond images:
 - Precomputation & look-up tables
 - Normal maps, reflection maps, shadow maps
- Procedural textures

