

---

# Computer Graphics

## Textures, part II

Konstantin Tretyakov

kt@ut.ee



---

Oct 30, 2013

# Standard Graphics Pipeline

---

Vertex  
transform

**Determine clip-space position of a triangle**

Culling and  
clipping

**Determine whether the triangle is visible**

Rasterization

**Determine all pixels belonging to the triangle**

Fragment  
shading

**For each pixel, determine its color**

Visibility tests  
& blending

**Draw pixel (*if needed*)**



# Quiz

---

- Formulate the Blinn model:
  - One light source, one color component
  - Specular exponent 10
  - Quadratic attenuation
  - Spotlight exponent 15
  - Exponential fog with coefficient 2
  - Fog color black



# Quiz

---

- Formulate the Blinn model:
  - One light source, one color component
  - Specular exponent 10
  - Quadratic attenuation
  - Spotlight exponent 15
  - Exponential fog with coefficient 2
  - Fog color black

$$I_A M_A +$$



# Quiz

---

- Formulate the Blinn model:
  - One light source, one color component
  - Specular exponent 10
  - Quadratic attenuation
  - Spotlight exponent 15
  - Exponential fog with coefficient 2
  - Fog color black

$$I_A M_A + (I_D M_D \cdot (\mathbf{n}^T \mathbf{l})_+ +$$



# Quiz

---

- Formulate the **Blinn** model:
  - One light source, one color component
  - **Specular exponent 10**
  - Quadratic attenuation
  - Spotlight exponent 15
  - Exponential fog with coefficient 2
  - Fog color black

$$I_A M_A + \left( I_D M_D \cdot (\mathbf{n}^T \mathbf{l})_+ + I_S M_S \cdot (\mathbf{n}^T \mathbf{h})_+^{10} \right)$$



# Quiz

---

- Formulate the **Blinn** model:
  - One light source, one color component
  - **Specular exponent 10**
  - **Quadratic attenuation**
  - Spotlight exponent 15
  - Exponential fog with coefficient 2
  - Fog color black

$$I_A M_A + \left( I_D M_D \cdot (\mathbf{n}^T \mathbf{l})_+ + I_S M_S \cdot (\mathbf{n}^T \mathbf{h})_+^{10} \right) \frac{1}{d^2}$$




# Quiz

---

- Formulate the **Blinn** model:
  - One light source, one color component
  - **Specular exponent 10**
  - **Quadratic attenuation**
  - **Spotlight exponent 15**
  - Exponential fog with coefficient 2
  - Fog color black

$$I_A M_A + \left( I_D M_D \cdot (\mathbf{n}^T \mathbf{l})_+ + I_S M_S \cdot (\mathbf{n}^T \mathbf{h})_+^{10} \right) \frac{1}{d^2} (-\mathbf{l}^T \mathbf{s})^{15}$$






# Quiz

---

- Formulate the **Blinn** model:
  - One light source, one color component
  - **Specular exponent 10**
  - **Quadratic attenuation**
  - **Spotlight exponent 15**
  - **Exponential fog with coefficient 2**
  - Fog color black

$$\text{mix}( I_A M_A + (I_D M_D \cdot (\mathbf{n}^T \mathbf{l})_+ + I_S M_S \cdot (\mathbf{n}^T \mathbf{h})_+^{10}) \frac{1}{d^2} (-\mathbf{l}^T \mathbf{s})^{15}, \\ \text{black}, \\ \exp(-2|z|) )$$




# Quiz

---

- Formulate the **Blinn** model:
  - One light source, one color component
  - **Specular exponent 10**
  - **Quadratic attenuation**
  - **Spotlight exponent 15**
  - **Exponential fog with coefficient 2**
  - Fog color black


$$\exp(-2|z|) \cdot \left( I_A M_A + \left( I_D M_D \cdot (\mathbf{n}^T \mathbf{l})_+ + I_S M_S \cdot (\mathbf{n}^T \mathbf{h})_+^{10} \right) \frac{1}{d^2} (-\mathbf{l}^T \mathbf{s})^{15} \right)$$



# In the previous lecture

---

- We can specify color of every pixel of a triangle by mapping it from a *texture image*.



# In the previous lecture

---

- How attribute interpolation is *actually* done
- Texture filtering
  - Bilinear, mipmapping, anisotropic
- Textures & OpenGL
- Textures beyond images:
  - Precomputation & look-up tables
  - Normal maps, environment maps, shadow maps
- Procedural textures



# Quiz

---

- This is a 2x2 pixel checkerboard texture spread on a rectangle. What filtering setting is used?



# Quiz

---

- This is a 2x2 pixel checkerboard texture spread on a rectangle. What filtering setting is used?





```
glTexParameteri(GL_TEXTURE_2D,  
                 GL_TEXTURE_MAG_FILTER,  
                 GL_LINEAR)
```



# Quiz

---


- What filtering setting is used here?



# Quiz

---

- What filtering setting is used here?





# Next

---

- How attribute interpolation is *actually* done
- Texture filtering
  - Bilinear, mipmapping, anisotropic
- Textures & OpenGL
- Textures beyond images:
  - Precomputation & look-up tables
  - Normal maps, environment maps, shadow maps
- Procedural textures



# Textures as look-up tables

---


```
uniform sampler1D complexFunction;  
  
void main() {  
    ...  
    float light_attenuation =  
        texture1D(complexFunction, angle).x;  
    ...  
}
```



# Textures as look-up tables

---

```
uniform sampler3D specularIntensity;  
  
void main() {  
    ...  
    float specular_term =  
        texture3D(specularIntensity,  
                  reflectionDirection).x;  
    ...  
}
```



# Textures as look-up tables

---

```
uniform sampler2D randomness;


void main() {
    ...
    float effect =
        texture2D(randomness, gl_FragCoord.xy).x;
    ...
}
```





# Normal mapping

---

Use a texture to specify (or modify) the *normal* at each point of the surface.



# Normal mapping




# Bump-mapping

---

- Although it is possible to store the actual  $(x,y,z)$  normal direction for each texel, it is more common to use the texture to store a single number per texel – the *height* above the corresponding point.

Texture values (height)




# Bump-mapping

---

- The normal direction can be inferred from the height map via discrete differentiation:

$$n_x(uv) \propto h(uv - (\varepsilon, 0)) - h(uv)$$





# Bump-mapping

---

- The normal direction can be inferred from the height map via discrete differentiation:

$$\mathbf{n} \propto \begin{pmatrix} h(\mathbf{uv} - (\varepsilon, 0)) - h(\mathbf{uv}) \\ h(\mathbf{uv} - (0, \varepsilon, )) - h(\mathbf{uv}) \\ 1 \end{pmatrix}$$



# Bump-mapping

---

- The normal direction can be inferred from the height map via discrete differentiation:

$$\mathbf{n} \propto \begin{pmatrix} h(\mathbf{uv} - (\varepsilon, 0)) - h(\mathbf{uv}) \\ h(\mathbf{uv} - (0, \varepsilon, )) - h(\mathbf{uv}) \\ 1 \end{pmatrix}$$

Note that those are normal coordinates in the polygon-local coordinate system.



# Bump-mapping

---

- The normal direction can be inferred from the height map via discrete differentiation:

$$\mathbf{n} \propto \begin{pmatrix} h(\mathbf{uv} - (\varepsilon, 0)) - h(\mathbf{uv}) \\ h(\mathbf{uv} - (0, \varepsilon, )) - h(\mathbf{uv}) \\ 1 \end{pmatrix}$$

The normal's object coordinates are then

$$\mathbf{Mn}$$

Where  $\mathbf{M} = (\mathbf{u}, \mathbf{v}, \mathbf{w})$  provides the “up”, “right” and “out” directions wrt polygon



# Reflections

---

Suppose we are looking at some point of a highly reflective object



# Reflections

---

Suppose we are looking at some point of a highly reflective object



We should see the reflection of whatever is in that direction



# Reflections

---

Suppose we are looking at some point of a highly reflective object



We could in principle trace this ray until we figure out what is there in that direction.



# Environment mapping


---

Use a texture to store what is seen from a point for each possible direction.

Suppose we are looking at some point of a highly reflective object



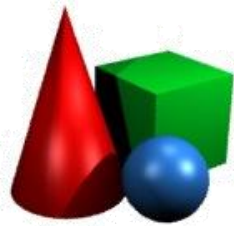

No need to trace the ray, if we have the necessary value precomputed.



# Environment mapping

---

Assume that this mapping is **the same for all points**, because the reflections are “distant” (hence “environment” mapping).





# Environment map

---

Environment map – a texture that stores for any possible direction the color of a distant point in that direction.



# Environment map

---

Environment map – a texture that stores for any possible direction the color of a distant point in that direction.

Sphere map

Cube map

HEALpix map


... or any cartographic projection



# Sphere map

---

An orthogonal view of a perfectly reflective sphere shows reflections for all possible directions.



# Sphere map


---

So now for any reflection direction we simply have to find the point on this image, where the sphere had the same reflection vector.



# Sphere mapping

---




Say we are shading this pixel  
on a reflective polygon



# Sphere mapping

---



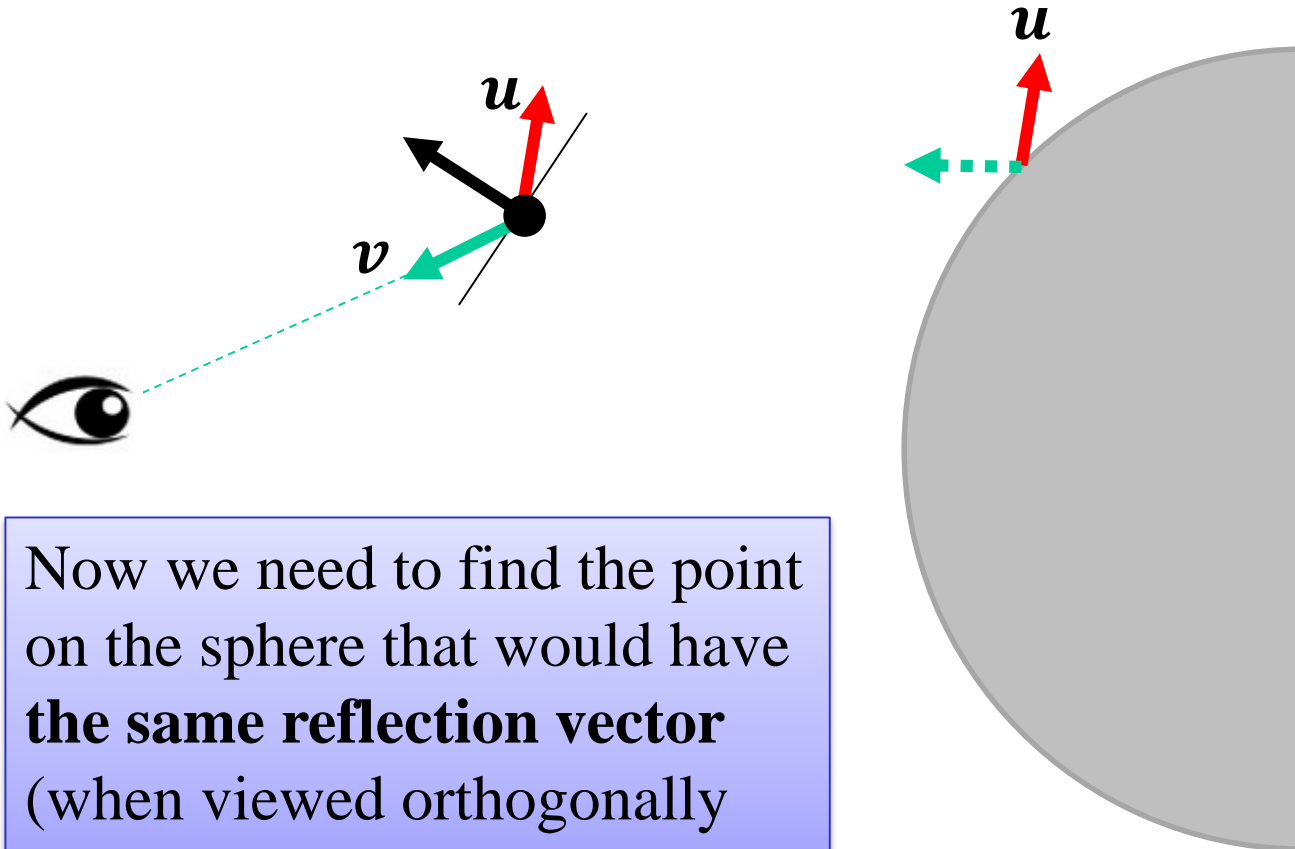
Say we are shading this pixel on a reflective polygon

Start by computing the viewer direction vector and its reflection against the normal



# Sphere mapping

---




Now we need to find the point on the sphere that would have **the same reflection vector** (when viewed orthogonally straight)



# Sphere mapping

---



1. Compute the normal to the sphere at that point:


$$n_s = \text{normalize}(u + (0,0,1))$$





# Sphere mapping

---




2. Find the actual point on the sphere

$$\mathbf{p} = (0.5, 0.5, 0) + 0.5\mathbf{n}_s$$



# Sphere mapping

---



2. Find the actual point on the sphere

$$\mathbf{p} = (0.5, 0.5, 0) + 0.5\mathbf{n}_s$$

3. Sample the sphere map texture at  $(p_x, p_y)$




# Sphere map

---

OpenGL can do it  
for us via  
automated texture  
coordinate  
generation:

```
glTexGeni(GL_S, GL_TEXTURE_GEN_MODE, GL_SPHERE_MAP);  
glTexGeni(GL_S, GL_TEXTURE_GEN_MODE, GL_SPHERE_MAP);
```



# See e.g. glTexGeni manual

---

If *pname* is GL\_SPHERE\_MAP and *coord* is either GL\_S or GL\_T, *s* and *t* texture coordinates are generated as follows. Let *u* be the unit vector pointing from the origin to the polygon vertex (in eye coordinates). Let *n'* be the current normal, after transformation to eye coordinates. Let *f* = (*f<sub>x</sub>* () *f<sub>y</sub>* () *f<sub>z</sub>*)<sup>T</sup> be the reflection vector such that

$$\mathbf{f} = \mathbf{u} - 2\mathbf{n}' \mathbf{n}'^T \mathbf{u}$$

Finally, let

$$m = 2 \sqrt{f_x^2 + f_y^2 + (f_z + 1)^2}$$

Then the values assigned to the *s* and *t* texture coordinates are

$$s = \frac{f_x}{m} + \frac{1}{2}$$

$$t = \frac{f_y}{m} + \frac{1}{2}$$



# Cube map

---

- Sphere map only works well for a fixed viewer direction, as the “sides” of the sphere map are heavily undersampled.



# Cube map

---

- Sphere map only works well for a fixed viewer direction, as the “sides” of the sphere map are heavily undersampled.
- **Cube map** is a set of 6 images, forming the “insides” of a cube. Now for the point in the center of the cube we can easily compute what is seen in any direction.
- This lets us compute reflections for any viewpoint with equal accuracy.



# Cube map

---



# Cube map

```
glEnable(GL_TEXTURE_CUBE_MAP);  
glBindTexture(GL_TEXTURE_CUBE_MAP, textureHandle);  
glTexImage2D(GL_TEXTURE_CUBE_MAP_POSITIVE_X, ...);  
glTexImage2D(GL_TEXTURE_CUBE_MAP_NEGATIVE_X, ...);  
glTexImage2D(GL_TEXTURE_CUBE_MAP_POSITIVE_Y, ...);  
glTexImage2D(GL_TEXTURE_CUBE_MAP_NEGATIVE_Y, ...);  
glTexImage2D(GL_TEXTURE_CUBE_MAP_POSITIVE_Z, ...);  
glTexImage2D(GL_TEXTURE_CUBE_MAP_NEGATIVE_Z, ...);  
  
glTexGeni(GL_S, GL_TEXTURE_GEN_MODE, GL_REFLECTION_MAP);  
glTexGeni(GL_T, GL_TEXTURE_GEN_MODE, GL_REFLECTION_MAP);
```

.. or use explicitly in the shader:

```
uniform samplerCube myCubeMap;  
...  
vec4 color = textureCube(myCubeMap, vec3(x, y, z));
```






# Shadow maps

---


Texture-based precomputation is a popular way of rendering shadows.



# Shadow maps

---


For a given light source we can render a texture, that stores for every light ray its length until it hits an object.



# Shadow maps

---


This texture can then be used when rendering the scene, to test for each fragment, whether it is occluded from the light source.



# Shadow maps

---

This texture can then be used when rendering the scene, to test for each fragment, whether it is occluded from the light source.



---


\* More formally, find its coordinates in the light source’s “projective frustum”



# Shadow maps

---

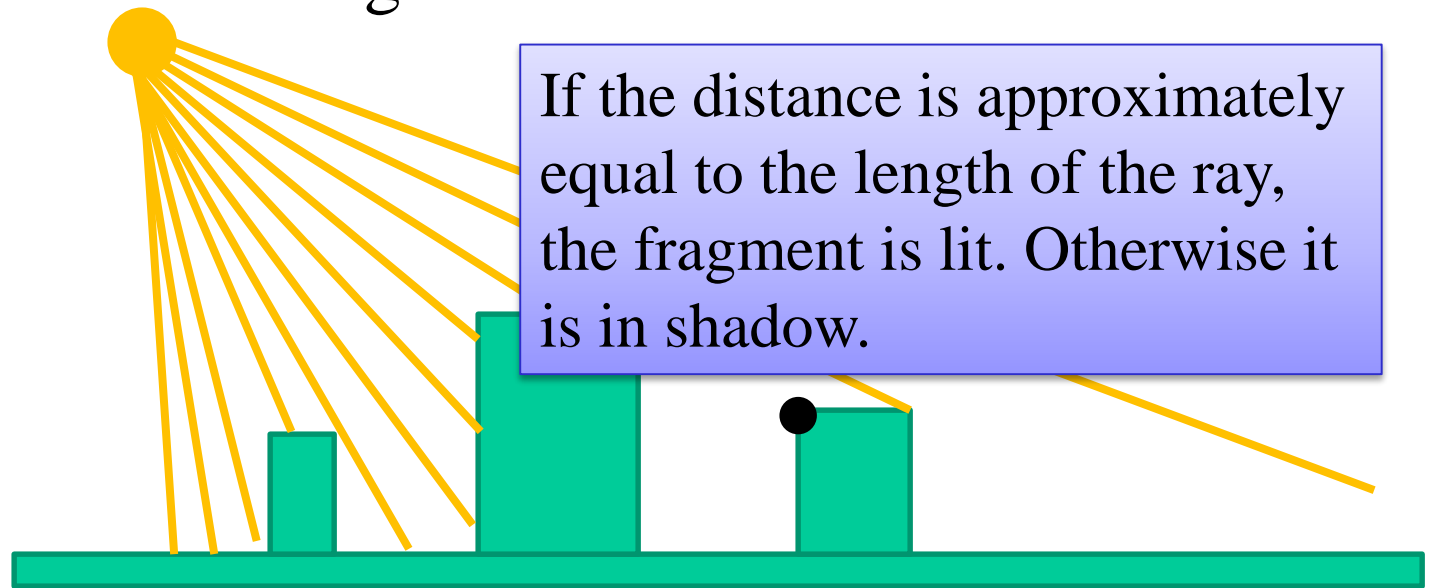
This texture can then be used when rendering the scene, to test for each fragment, whether it is occluded from the light source.



# Shadow maps

---


This texture can then be used when rendering the scene, to test for each fragment, whether it is occluded from the light source.



# Light maps


---

- Do not confuse the shadow map method with another straightforward approach to render shadows: simply precompute the texture for each object that takes lighting and shadows into account. This texture is called a **light map**.



# Light maps

---





# Texturing

---

- How attribute interpolation is *actually* done
- Texture filtering
  - Bilinear, mipmapping, anisotropic
- Textures & OpenGL
- Textures beyond images:
  - Precomputation & look-up tables
  - Normal maps, environment maps, shadow maps
- Procedural textures



# Texturing

---

- How attribute interpolation is *actually* done
- Texture filtering
  - Bilinear, mipmapping, anisotropic
- Textures & OpenGL
- Textures beyond images:
  - Precomputation & look-up tables
  - Normal maps, environment maps, shadow maps
- Procedural textures



# Procedural textures

---



- Often it makes sense to generate texture images procedurally.
- Moreover, you may sometimes do it on the fly rather than read from a file.
- I.e. a “texture” in this case is simply a function  $f(x)$  or  $f(x,y)$  or  $f(x,y,z)$ .



# Texture synthesis

---


- Texture synthesis is a large field in itself, here we shall only cover one most relevant algorithm: **Perlin fractal noise**.



# Basic 1D Perlin noise

---

- Fix a function  $f$ , that can produce a random-looking float for a fixed integer input.




```
float my_rand(int n) {  
    return cos(12345*n*n);  
}
```



# Basic 1D Perlin noise

---


- Use the value of  $f$  at each integer point  $n$  as the derivative of a linear function passing through  $(n, 0)$ :





# Basic 1D Perlin noise

- For points between integer positions, blend between corresponding functions:




```
t = fade(frac(x))  
return (1-t)*f1 + t*f2
```



# Basic 1D Perlin noise

---

- For points between integer positions, blend between corresponding functions:






# Basic 1D Perlin noise

---


- For points between integer positions, blend between corresponding functions:



# Basic 1D Perlin noise

---


- For points between integer positions, blend between corresponding functions:



# Basic 1D Perlin noise


---

- For points between integer positions, blend between corresponding functions:




# Fractal noise

- Add noise generated at several scales.




`2.0*perlin(0.5*x)`

+




`1.0*perlin(1.0*x)`

+



`0.2*perlin(2.0*x)`


=



# Multidimensional noise

---


- For 2D, 3D, etc, the logic is the same, but now we generate a **gradient vector** at each grid point.



# Multidimensional noise

---


- For 2D, 3D, etc, the logic is the same, but now we generate a **gradient vector** at each grid point.



# Multidimensional noise

---


- For 2D, 3D, etc, the logic is the same, but now we generate a **gradient vector** at each grid point.




# Fractal multidimensional Perlin noise

---

- Same idea for generating fractal noise – add together noise at several scales:



```
2.0*perlin2d(0.3*x) + 1.0*perlin2d(x) +  
0.3*perlin2d(2*x) + 0.1*perlin2d(4*x)
```





# Texturing

---

- How attribute interpolation is *actually* done
- Texture filtering
  - Bilinear, mipmapping, anisotropic
- Textures & OpenGL
- Textures beyond images:
  - Precomputation & look-up tables
  - Normal maps, environment maps, shadow maps
- Procedural textures



# Standard Graphics Pipeline

---

Vertex  
transform

**Determine clip-space position of a triangle**

Culling and  
clipping

**Determine whether the triangle is visible**

Rasterization

**Determine all pixels belonging to the triangle**

Fragment  
shading

**For each pixel, determine its color**

Visibility tests  
& blending

**Draw pixel (*if needed*)**



# Interlude: Shadows

---



- Light maps
- Shadow maps
- Projective shadows
- Stencil shadows



# Projective shadows



---

- “Flatten” objects into their shadows via a projection transform from the light source onto a plane.



# Projective shadows

- ```
draw_object();  
set_shadow_color();  
glPushMatrix();  
    shadowProjectionMatrix(light, plane);  
    draw_object();  
glPopMatrix();
```





# Stencil Shadows

---



# Stencil Shadows

---



# Stencil Shadows

---

- Idea:
  - For each pixel, check whether it is within a shadow volume.
  - This can be done by counting
    - ▶ How many front-facing shadow volume polygons are in front of the pixel
    - ▶ How many back-facing shadow volume polygons are in front of the pixel
  - If the numbers do not match, the pixel is in shadow





# Stencil Shadows

---

- Implementation:
  - First render the scene as if it was all in shadow.
  - Render the front-facing shadow volume polygons, incrementing the stencil buffer for each pixel where the depth test passes.
  - Render back-facing shadow volume polygons, decrementing the stencil buffer where the depth test passes.
  - Now whenever stencil buffer is 0, render the scene as fully lit.

