

---

# Computer Graphics

Raycasting. Raytracing. Raymarching.

Konstantin Tretyakov  
kt@ut.ee



---


Nov 13, 2013

# Next

---

- Raycasting
- Raytracing
- Raymarching / Sphere tracing
- Rendering equation solvers
  - Radiosity, Photon mapping, Path tracing





# The “Emission theory”

---

- Ancient greeks (up to Aristotle and even later) believed that you see by using beams emanating from your eyes.
- “Winer et al. (2002) have found recent evidence that as many as 50% of American college students believe in emission theory.[\[3\]](#)”



# The “Emission theory”

---


- Mathematically, the two theories are equivalent in many respects, because light propagation rules are mostly direction-independent.
- Consequently, we can use the “Emission” approach in rendering. It is called raycasting.



# Graphics Pipeline vs Raycasting

---


- **Graphics pipeline:**
  - For each object, find its location on screen (and then color each pixel appropriately)



# Graphics Pipeline vs Raycasting

---


- **Graphics pipeline:**
  - For each object, find its location on screen (and then color each pixel appropriately)



# Graphics Pipeline vs Raycasting

---

- **Graphics pipeline:**
  - For each object, find its location on screen (and then color each pixel appropriately)






# Graphics Pipeline vs Raycasting

---


- **Graphics pipeline:**
  - For each object, find its location on screen (and then color each pixel appropriately)



# Graphics Pipeline vs Raycasting

---


- **Raycasting:**
  - For each **pixel**, find the corresponding **object** (and then color each pixel appropriately)



# Graphics Pipeline vs Raycasting

---


- **Raycasting:**
  - For each **pixel**, find the corresponding **object** (and then color each pixel appropriately)



# Graphics Pipeline vs Raycasting

---


- **Raycasting:**
  - For each **pixel**, find the corresponding **object** (and then color each pixel appropriately)



# Example: Plane

---


Ray-plane intersection



# Example: Plane

---


Ray-plane intersection



# Example: Plane

---

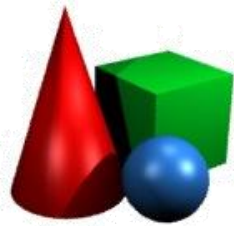

Ray-plane intersection



# Example: Plane

---

Ray-plane intersection






# Example: Plane

---


Ray-plane intersection



# Example: Plane

---


Ray-plane intersection



# Example: Plane

---


Ray-plane intersection



# Example: Plane

---

Ray-plane intersection



(assuming  $\|\mathbf{n}\| = 1$ )

(assuming  $\|\mathbf{v}\| = 1$ )



# Example: Plane

---

```
float raycast_plane(vec3 from, vec3 dir,
                    vec3 n,    float b) {
    if (dot(n, dir) == 0)
        return -1;
    else
        return -(dot(n, from) + b) / dot(n, dir);
}
```



# Example: Plane

---

```
float raycast_plane(vec3 from, vec3 dir,  
                    vec3 n,    float b) {  
    if (dot(n, dir) == 0)  
        return -1;  
    else  
        return -(dot(n, from) + b) / dot(n, dir);  
}
```

Separate ray-surface intersection routine can be derived for every type of surface.



# Quiz

---

- Suppose we have  $n$  objects, each with screen area  $\sim 1000$  pixels. Let the total screen area be  $800 \times 600 = 480\,000$  pixels.
- Complexity of (naïve) graphics pipeline:
  - ?
- Complexity of (naïve) raycasting:
  - ?



# Quiz

---

- Suppose we have  $n$  objects, each with screen area  $\sim 1000$  pixels. Let the total screen area be  $800 \times 600 = 480\,000$  pixels.
- Complexity of (naïve) graphics pipeline:
  - $1000\,n$
- Complexity of (naïve) raycasting:
  - ?





# Quiz

---

- Suppose we have  $n$  objects, each with screen area  $\sim 1000$  pixels. Let the total screen area be  $800 \times 600 = 480\,000$  pixels.
- Complexity of (naïve) graphics pipeline:
  - $1000\,n$
- Complexity of (naïve) raycasting:
  - $480\,000\,n$



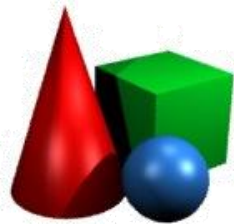
# Quiz

---

- Suppose we have  $n$  objects, each with screen area  $\sim 1000$  pixels. Let the total

In the worst case (each object occupies the whole screen) the complexities are equal.

- $1000 n$
- Complexity of (naïve) raycasting:
  - $480\,000 n$



# Graphics Pipeline vs Raycasting

---

- Theoretically, the complexity of naïve raycasting is not worse than that of the graphics pipeline.



# Graphics Pipeline vs Raycasting

---


- Theoretically, the complexity of naïve raycasting is not worse than that of the graphics pipeline.
- In fact, in the old days, properly-done (1D) raycasting was more efficient.



# Graphics Pipeline vs Raycasting

---

- However,
  - Raycasting parallelizes over **outputs** (pixels) and uses **full scene data for each pixel**.
  - GP parallelizes over **inputs** (splits input into triangles) which makes memory access of each subtask *local*.
- For large scenes the latter approach fits the architecture of modern hardware much better (shared memory access is slow).



# Benefits of Raycasting

---

- On the other hand, having full access to the scene while rendering lets you implement **global illumination** effects much more naturally than in the standard graphics pipeline.



# Global vs Local illumination


---

- **Local illumination model:**
  - Use only the information about light sources in lighting computations. Other objects in the scene do not matter.
- **Global illumination model:**
  - Take into account light scattering and occlusion due to **all** objects in the scene:
    - ▶ Shadows, reflection, refraction, ambient scattering, ambient occlusion



# Global vs Local illumination

---



---

[http://en.wikipedia.org/wiki/File:Recursive\\_raytrace\\_of\\_a\\_sphere.png](http://en.wikipedia.org/wiki/File:Recursive_raytrace_of_a_sphere.png)





# Raytracing


---

- **Raytracing** :=
  - Basic raycasting
  - + Shadows
  - + Reflections
  - + Refractions
  - + Recursion
  - + Backwards raytracing
  - + Stochastic raytracing



# Shadows via Raycasting


---



# Shadows via Raycasting

---

For each point, cast a ray towards the light source (“shadow ray”).




# Shadows via Raycasting

---

For each point, cast a ray towards the light source (“shadow ray”).

If the ray intersects some geometry before reaching the light source, the point is in shadow.



# Reflections via Raycasting

---

- For each point, continue the ray in a **reflected** direction, and combine the local lighting of the point with whatever is hit by the reflected ray.

$$C_{local\ model} + k_{refl}C_{refl}$$

(where  $k_{refl}$  is the “reflectivity coefficient”)



# Refractions via Raycasting

---


- For each point, continue the ray in a **refracted** direction, and combine the local lighting of the point with whatever is hit by the reflected ray + refracted ray

$$C_{local\ model} + k_{refl}C_{refl} + k_{refr}C_{refr}$$



# Recursive Raytracing

---



---

Illustration © Mark Fishel



# Faster Raytracing

---

- **BVHs & BSPs:**
  - Quiz: What were those?





# Faster Raytracing

---

- **BVHs & BSPs:**
  - Bounding Volume Hierarchies allow to quickly cull away irrelevant portions of the scene.  
Binary Space Partitioning trees let you seek for ray intersections in a front-to-back order.



# Faster Raytracing

---

- “Object z-buffer”:
  - Pre-render the scene using the Graphics pipeline, storing for each pixel the id of the topmost object rendered at that pixel.
  - This effectively replaces the first raycasting step with a graphics pipeline, making it faster.



# Faster Raytracing

---

- **Adaptive recursion depth:**
  - Keep track, for each ray, what will be its final contribution to the pixel it originated from.
  - E.g. if after hitting a surface we reflect with reflection coefficient 0.2, then refract with coefficient 0.1, we may stop as further computations won't contribute more than 0.02.



# Faster Raytracing

---


- **Shadow mapping:**
  - Instead of casting each “shadow ray”, use a precomputed depth buffer like in shadow mapping.



# Refracted lighting problem

---



**Refracted** light from light sources is not taken into account



# Refracted lighting problem

---

**Refracted** light from light sources is not taken into account



# Backwards Raytracing

---

- **Solution:**
  - Split geometry into patches.
  - Make a raytracing pass with rays originating from the light sources
  - For each patch store its lighting data.
  - Use this data in the normal forward pass.



# Aliasing

---

- Aliasing can be a considerable problem for raytracing implementations.
- Solutions:
  - Cone tracing, Beam tracing, Pencil tracing
  - **Stochastic raytracing**





# Stochastic raytracing


---

- Recall that a common way of dealing with aliasing artifacts is by using a convolution with a filter.
- In layman's terms: send multiple rays for each pixel, sampling slightly different directions, and average the results.



# Basic antialiasing

---



# Multiple rays per pixel


---

- For each ray we can vary many parameters:
  - Ray direction (basic antialiasing)
  - Ray origin (depth of field)
  - Object positions (motion blur)
  - Light source positions (soft shadows)
  - Ray light frequency (true refractions)



# Depth of Field

---



# True refraction

---

One ray per light frequency



# Combined stochastic rays

---

- Suppose we want to do
  - 4x4-grid jittered antialiasing
  - + 4x4-grid jittered DOF.
  - + 10-frame motion blur
- How many rays per pixel should we send?



# Combined stochastic rays

---

- Suppose we want to do
  - 4x4-grid jittered antialiasing
  - + 4x4-grid jittered DOF.
  - + 10-frame motion blur
- We do not need to send  $16 \times 16 \times 10 = 2560$  rays for all possible combinations!
- Choosing, say, 4x4 random parameter combinations can suffice.



# Outline

---

- Raycasting
- Raytracing
- Raymarching / Sphere tracing
- Rendering equation solvers
  - Radiosity, Photon mapping, Path tracing





# Outline

---


- Raycasting
  - Raytracing
  - Raymarching / Sphere tracing
- 
- Rendering equation solvers
    - Radiosity, Photon mapping, Path tracing



# Problems with Raytracing

---


- Finding ray intersections with complex shapes can be hard, and not always analytically possible.



# Raymarching

---


- Idea: use implicit surface representation:
  - $f(\mathbf{x}) = 0$ , for points  $\mathbf{x}$  on the surface.
  - $f(\mathbf{x}) > 0$ , for points outside the shape.



# Raymarching

---


- Now we can “walk” along the ray in small steps, seeking for the point where  $f(x) = 0$



# Sphere tracing


---

- Even better if  $f(\mathbf{x})$  is a **lower bound on the distance** to any shape in the scene. Then we can safely use step size  $f(\mathbf{x})$  along the ray.



# Sphere tracing


---



# Sphere tracing

---


We can safely proceed along the ray  
by at least 4.9



# Sphere tracing

---

Evaluate  $f$  at the current point






# Sphere tracing

---


... and proceed again



# Sphere tracing

---

If  $f(x)$  corresponds to an actual distance (or a lower bound with some additional criteria), the algorithm will converge towards the true intersection.



# Sphere tracing / Raymarching

---

```
float raymarch(vec3 from, vec3 dir, function f)
{
    vec3 x = from;
    int steps = MAX_STEPS;
    while (f(x) > 0.0001 && steps-- > 0)
        x += f(x)*dir;
    return length(x - from);
}
```



# Distance fields

---

- To do raymarching we have to model the whole scene in the form of a single *distance field function*  $f(\mathbf{x})$ .
- It turns out that distance fields let you model quite complicated shapes rather naturally.



# Quiz

---

- Devise a distance function for a sphere of radius 1 centered at  $(0, 0, 0)$ .



# Quiz

---

- Devise a distance function for a sphere of radius 1 centered at (0, 0, 0).

```
float sphere(vec3 x) {  
    return length(x) - 1;  
}
```




### Sphere - signed

```
float sdSphere( vec3 p, float s )
{
    return length(p)-s;
}
```




### Box - unsigned

```
float udBox( vec3 p, vec3 b )
{
    return length(max(abs(p)-b,0.0));
}
```




### Round Box - unsigned

```
float udRoundBox( vec3 p, vec3 b, float r )
{
    return length(max(abs(p)-b,0.0))-r;
}
```



### Box - signed

```
float sdBox( vec3 p, vec3 b )
{
    vec3 d = abs(p) - b;
    return min(max(d.x,max(d.y,d.z)),0.0) +
        length(max(d,0.0));
}
```




### Torus - signed

```
float sdTorus( vec3 p, vec2 t )
{
    vec2 q = vec2(length(p.xs)-t.x,p.y);
    return length(q)-t.y;
}
```



### Cylinder - signed

```
float sdCylinder( vec3 p, vec3 c )
{
    return length(p.xs-c.xy)-c.z;
}
```



### Cone - signed

```
float sdCone( vec3 p, vec2 c )
{
    // c must be normalised
    float q = length(p.xy);
    return dot(c,vec2(q,p.z));
}
```



### Plane - signed

```
float sdPlane( vec3 p, vec4 n )
{
    // n must be normalised
    return dot(p,n.xyz) + n.w;
}
```



# Quiz

---

- Given two distance functions corresponding to some shapes, devise a distance function that corresponds to the union of these shapes:





# Quiz

---

- Given two distance functions corresponding to some shapes, devise a distance function that corresponds to the union of these shapes:

```
float union(vec3 x) {  
    return min(f1(x), f2(x));  
}
```




# Constructive Solid Geometry

---

## Union

```
float opU( float d1, float d2 )
{
    return min(d1,d2);
}
```




## Intersection

```
float opI( float d1, float d2 )
{
    return max(d1,d2);
}
```



## Subtraction

```
float opS( float d1, float d2 )
{
    return max(-d1,d2);
}
```



## Blend

```
float opBlend( vec3 p )
{
    float d1 = primitiveA(p);
    float d2 = primitiveB(p);
    return smin( d1, d2 );
}
```



# Quiz

---

- Given a distance function describing a shape centered at  $(0, 0, 0)$ . Devise a distance function for the same shape, centered at  $(1, 1, 1)$ :



# Quiz

---

- Given a distance function describing a shape centered at  $(0, 0, 0)$ . Devise a distance function for the same shape, centered at  $(1, 1, 1)$ :

```
float translate(vec3 x) {  
    return f(x - vec3(1, 1, 1))  
}
```



# Arbitrary transformations

## Rotation/Translation

```
vec3 opTx( vec3 p, mat4 m )
{
    vec3 q = invert(m)*p;
    return primitive(q);
}
```



## Scale

```
float opScale( vec3 p, float s )
{
    return primitive(p/s)*s;
}
```




## Repetition

```
float opRep( vec3 p, vec3 c )
{
    vec3 q = mod(p,c)-0.5*c;
    return primitve( q );
}
```



## Displacement

```
float opDisplace( vec3 p )
{
    float d1 = primitive(p);
    float d2 = displacement(p);
    return d1+d2;
}
```




## Twist

```
float opTwist( vec3 p )
{
    float c = cos(20.0*p.y);
    float s = sin(20.0*p.y);
    mat2 m = mat2(c,-s,s,c);
    vec3 q = vec3(m*p.xz,p.y);
    return primitive(q);
}
```



## Cheap Bend

```
float opCheapBend( vec3 p )
{
    float c = cos(20.0*p.y);
    float s = sin(20.0*p.y);
    mat2 m = mat2(c,-s,s,c);
    vec3 q = vec3(m*p.xy,p.z);
    return primitive(q);
}
```





---

© Iñigo Quílez, <http://www.iquilezles.org/>



# Next time

---

- Raycasting
- Raytracing
- Raymarching / Sphere tracing
- Rendering equation solvers
  - Radiosity, Photon mapping, Path tracing

